



This is the Thesis Title

*A thesis submitted in partial fulfillment of the requirements of the B.Sc.
degree in Electrical and Electronic Engineering*

Submitted by:

Student Name (Index)

Student Name (Index)

Supervised by:

Supervisor name

To

Department of Electrical and Electronic Engineering

Faculty of Engineering

University of Khartoum

September 2024

Declaration of Authorship

You can check the official EEE department thesis guidelines from the main menu **"thesis Guidelines .pdf "**

Dedication

Acknowledgments

Abstract

Retrieval-Augmented Generation (RAG) systems depend on effective information retrieval to ground Large Language Model outputs in external knowledge, yet Arabic information retrieval faces severe challenges stemming from morphological richness, orthographic variation, and vocabulary mismatch between queries and documents. Query enhancement—the modification of a user’s query before retrieval—offers a modular, retriever-agnostic intervention to address these challenges. However, existing query enhancement techniques were developed for English using proprietary models with 175 billion or more parameters, and their transferability to Arabic using small, open-source models has not been systematically investigated.

This thesis investigates the extent to which small open-source LLMs can improve Arabic information retrieval through query enhancement, and identifies the model characteristics that determine effectiveness. Experiments were conducted on the MIRACL Arabic benchmark, comprising 2,896 queries and 2.06 million passages, using mDPR for dense retrieval and BM25S for sparse retrieval. Error analysis on the baseline identified a 34% query failure rate and the pronounced underperformance of very short queries, motivating the adaptation of the Query2Doc technique for Arabic zero-shot application. Ten open-source LLMs with 2–8 billion parameters were evaluated using a standardised protocol, with all experiments executed on Google Colab GPUs.

Query enhancement improved dense retrieval for all nine viable models, with NDCG@10 improvements ranging from +3.7% to +23.5%. Aya Expanse 8B achieved the highest overall improvement (+23.5% dense, +9.2% BM25), while Jais-2 8B was the most effective model for sparse retrieval (+10.8%). Zero-shot Arabic query enhancement with a 3-billion-parameter model exceeded the original Query2Doc paper’s few-shot English results obtained with a 175-billion-parameter model. However, BM25 retrieval was degraded by six of nine

models due to term dilution, demonstrating that query enhancement strategies must be adapted to the retrieval paradigm. Model size was positively correlated with dense retrieval improvement, and training data diversity was found to be a more significant predictor of query expansion quality than Arabic-specific benchmark scores.

The proposed corpus-steered pipeline, combined with hybrid BM25+Dense fusion, achieved 0.7137 NDCG@10 on the MIRACL Arabic benchmark—a 54.5% improvement over the BM25 baseline and a 13.9% improvement over a strong no-QE hybrid system.

المستخلص

تعتمد أنظمة التوليد المعزّز بالاسترجاع على فعالية استرجاع المعلومات لربط مخرجات النماذج اللغوية الكبيرة بالمعرفة الخارجية، غير أن استرجاع المعلومات باللغة العربية يواجه تحديات جوهرية ناجمة عن الثراء الصرفي والتباين الإملائي وعدم التطابق المعجمي بين الاستعلامات والمستندات. يُمثّل تحسين الاستعلام---أي تعديل استعلام المستخدم قبل عملية الاسترجاع---تدخلاً نمطياً مستقلاً عن نوع المسترجع لمعالجة هذه التحديات. إلا أن تقنيات تحسين الاستعلام الحالية طُوّرت للغة الإنجليزية باستخدام نماذج مملوكة تتجاوز 175 مليار معامل، ولم تُدرّس بشكل منهجي إمكانية نقلها إلى اللغة العربية باستخدام نماذج صغيرة مفتوحة المصدر.

تبحث هذه الأطروحة في مدى قدرة النماذج اللغوية الكبيرة الصغيرة مفتوحة المصدر على تحسين استرجاع المعلومات العربية من خلال تحسين الاستعلام، وتحديد خصائص النموذج التي تؤثر في الفعالية. أُجريت التجارب على معيار MIRACL العربي المكوّن من 2,896 استعلاماً و 06.2 مليون مقطع نصي، باستخدام mDPR للاسترجاع الكثيف و BM25S للاسترجاع المتناثر. كشف تحليل الأخطاء على خط الأساس عن معدل فشل بلغ 34% وضعف ملحوظ في أداء الاستعلامات القصيرة جداً، مما حفّز تكييف تقنية Query2Doc للتطبيق العربي بدون أمثلة تدريبية. قُيِّمت عشرة نماذج لغوية كبيرة مفتوحة المصدر بأحجام تتراوح بين 2 و 8 مليارات معامل باستخدام بروتوكول تقييم موحد، مع تنفيذ جميع التجارب على وحدات معالجة الرسومات في Colab. Google

حقّق تحسين الاستعلام تحسناً في الاسترجاع الكثيف لجميع النماذج التسعة القابلة للاستخدام، بنسب تحسّن في NDCG@10 تراوحت بين 7%.3+ و 5%.23+ حقّق نموذج Expanse Aya 8B أعلى تحسّن إجمالي (+5%.23 كثيف، +2%.9 BM25) بينما كان 8B Jais-2 الأكثر فعالية للاسترجاع المتناثر (+8%.10). تجاوز تحسين الاستعلام العربي بدون أمثلة تدريبية باستخدام نموذج بحجم 3 مليارات معامل النتائج الأصلية لورقة Query2Doc التي استخدمت نموذجاً بحجم 175 مليار معامل مع أمثلة تدريبية. غير أن أداء BM25 تراجع مع ستة من تسعة نماذج بسبب تخفيف المصطلحات، مما يؤكد ضرورة تكييف استراتيجيات تحسين الاستعلام وفقاً لنموذج الاسترجاع المستخدم. ارتبط حجم النموذج إيجابياً بتحسّن الاسترجاع الكثيف، وتبيّن أن تنوّع بيانات التدريب مؤثر أهم لجودة توسيع الاستعلام من درجات المعايير العربية المتخصصة.

حققت منظومة التوسيع الموجّه بالمدوّنة المقترحة، مقترنةً بالدمج الهجين بين BM25 والاسترجاع الكثيف، قيمة 7137.0 في $NDCG@10$ على معيار MIRACL العربي---بتحسّن نسبته 54.5%. مقارنةً بخط أساس BM25 وتحسّن نسبته 13.9% مقارنةً بنظام هجين قوي بدون تحسين للاستعلام.

List of Tables

2.1	Query-enhancement papers reviewed for this thesis, with method family, evaluation dataset, and headline result.	13
2.2	Arabic retrieval and question-answering datasets surveyed for benchmark selection. “Suitability” refers to fitness for studying query enhancement for retrieval.	18
2.3	Open-source LLMs evaluated for Arabic query enhancement in this thesis, ordered by parameter count.	20
2.4	Summary of language models used in the experiments.	26
3.1	LLM generation parameters for the initial Query2Doc experiment.	44
3.2	Model configurations for the Query2Doc comparison experiments. All models used <code>max_new_tokens</code> = 128 and <code>top_p</code> = 0.9 unless otherwise noted.	47
4.1	Baseline retrieval results on MIRACL Arabic dev set (2,896 queries). Bold marks the higher score between the two single-retriever baselines; the no-QE hybrid fusion result is included as the strongest non-enhancement reference (detailed in Section 4.7).	58
4.2	Performance segmentation of baseline queries by NDCG@10 score.	60
4.3	Retrieval performance by query length bucket.	61
4.4	Proportion of queries with at least one relevant passage at each retrieval depth.	63
4.5	Query2Doc results with dense retrieval (mDPR). Improvement is computed relative to the dense baseline.	64
4.6	Query2Doc results with BM25 retrieval (BM25S). Negative values indicate performance degradation.	65

4.7	Comparison between the original Query2Doc results (English, MS-MARCO) and this implementation (Arabic, MIRACL).	66
4.8	Dense retrieval (mDPR) results for all models, ranked by NDCG@10. The improvement column is relative to the no-enhancement dense baseline. .	67
4.9	BM25 retrieval (BM25S) results for all models, ranked by NDCG@10. The improvement column is relative to the no-enhancement BM25 baseline. Negative values indicate degradation.	69
4.10	Generational comparison within the Qwen model family.	72
4.11	Number of models improving over baseline, by retriever type.	73
4.12	BM25 nDCG@10 for nine models under query repetition. Columns $n \in \{1, 3, 5, 7, 10\}$ are fixed-count repetition (Query2Doc style); columns $\beta \in \{2, 4, 6\}$ are MuGI adaptive repetition where n is computed per-query as $n = \max(1, \lfloor d /(q \cdot \beta) \rfloor)$. MIRACL Arabic dev (2,896 queries). Best value per model in bold.	74
4.13	BM25 retrieval metrics at each model's best repetition configuration. Δ nDCG is the improvement from $n = 1$ (no repetition) to the best configuration. .	75
4.14	Hybrid retrieval fusion results on the MIRACL Arabic dev set (2,896 queries). CC = Convex Combination (min-max normalised per query, α controls Dense weight); RRF = Reciprocal Rank Fusion.	77
4.15	nDCG@10, Recall@10, Recall@100, and MRR for CSQE-enhanced retrieval on the MIRACL Arabic dev set. CSQE config: $k = 5$ first-pass, 2 corpus + 2 blind samples, $\alpha = 4$, Aya Expanse 8B.	79
4.16	BM25 retrieval metrics for corpus-only, blind-only, and combined CSQE expansion variants. All variants use $\alpha = 4$, Aya Expanse 8B, $k = 5$ first-pass. .	80
4.17	Effect of query repetition factor α on CSQE BM25 retrieval nDCG@10. .	80
4.18	Results for three CSQE application strategies in hybrid fusion. BM25-expanded: BM25 receives CSQE query, Dense receives original query. Dense-expanded: BM25 receives original query, Dense receives CSQE query. Both-expanded: both retrievers receive CSQE query.	81
4.19	BM25-expanded RRF ablation: contribution of corpus and blind expansion components when fused with the original Dense run.	82

4.20	Effect of query repetition factor α on the best BM25-expanded RRF system. All values use the same CSQE 2c+2b expansion, varying only the α prepended to the BM25-side query.	82
4.21	nDCG@10 improvements of BM25-expanded RRF over prior systems. . .	83
4.22	nDCG@10 progression from baseline to best system on the MIRACL Arabic dev set (2,896 queries). Δ cumulative measured from BM25S baseline.	84
4.23	Per-query comparison of the CSQE+Hybrid system (BM25-expanded RRF) against the Aya blind BM25 $n = 1$ baseline (0.5046 nDCG@10) on 2,896 MIRACL Arabic dev queries. The system's per-query-mean nDCG@10 is 0.6936; this is computed over individual queries and differs slightly from the corpus-level headline of 0.7137 reported in Table 4.22, which is produced by the official pooled evaluation.	85
4.24	System performance split by whether the BM25 first-pass retrieved a relevant document at rank 1 (top-1, qrel ≥ 1).	86
4.25	System performance decomposed by the number of words in the original MIRACL query, using the 1–3 / 4–8 / 9+ word buckets adopted throughout this thesis.	86
4.26	Representative big-win queries illustrating the corpus grounding effect. In all three cases, CSQE achieved nDCG@10 = 1.000 while the blind baseline achieved 0.000.	89
4.27	Classification of 367 regression queries (CSQE nDCG@10 < Blind by > 0.1).	89
4.28	Summary of all experiments grouped by phase: baselines, initial Query2Doc experiments, and expanded experiments covering query repetition, hybrid fusion, and corpus-steered query expansion. Bold values indicate the best result for each metric across all experiments.	92

List of Figures

2.1	RAG system architecture	9
3.1	Experimental pipeline overview	37
3.2	mDPR encoding flow	38
3.3	BM25S indexing flow	39
3.4	Query2Doc generation	43
3.5	Hybrid fusion: CC and RRF	52
3.6	CSQE pipeline	53
3.7	Best system architecture	56
4.1	Per-query NDCG@10 distribution	60
4.2	Failure cliff on baseline NDCG@10	61
4.3	NDCG@10 by query length	62
4.4	Recall@k curve	63
4.5	Dense retrieval NDCG@10 across models	68
4.6	Dense retrieval grouped multi-metric per model	68
4.7	Model size vs Query2Doc gain	71
4.8	BM25 query-repetition sweep across 9 models	76
4.9	Dense vs BM25 gain per model	76
4.10	Hybrid CC alpha sweep, all metrics	78
4.11	System progression to the best system	83
4.12	Per-query Delta NDCG@10 histogram	86
4.13	NDCG@10 by first-pass relevance	87
4.14	Per-bucket comparison by query length	88
4.15	Regression type breakdown	90

Contents

List of Tables	viii
List of Figures	ix
Contents	x
List of Abbreviations	xi
1 Introduction	1
1.1 Problem Definition	2
1.2 Objectives	3
1.3 Thesis Layout	4
2 Theoretical Background and Literature Review	6
2.1 Theoretical Background	6
2.1.1 Large Language Models and the Transformer Architecture	6
2.1.2 Retrieval-Augmented Generation	8
2.1.3 Information Retrieval Methods	9
2.1.4 Query Enhancement Techniques	11
2.1.5 Arabic Language Processing Challenges	13
2.2 Mathematical Models	14
2.2.1 BM25 Scoring Function	14
2.2.2 Dense Retrieval and Cosine Similarity	15
2.2.3 Hybrid Fusion Functions	15
2.2.4 Evaluation Metrics	16
2.3 Evaluation Dataset Selection	18
2.4 Models Used in Experiments	19
2.4.1 Arabic-Specialized Models	20

2.4.2	Multilingual Models	23
2.4.3	Retrieval Models	25
2.4.4	Model Comparison Summary	26
2.5	Related Work	27
2.5.1	Foundational Query Enhancement Research (2022–2023)	27
2.5.2	Modern LLM-Based Query Enhancement (2024–2025)	28
2.5.3	Arabic Information Retrieval and RAG	30
2.5.4	Research Gap	31
2.6	Chapter Summary	32
3	Methodology	34
3.1	Dataset and Experimental Setup	34
3.1.1	Dataset: MIRACL Arabic	34
3.1.2	Hardware and Software Environment	35
3.1.3	Evaluation Metrics	36
3.1.4	Experimental Pipeline Overview	36
3.2	Baseline Implementation	37
3.2.1	Dense Retrieval Baseline (mDPR)	37
3.2.2	Sparse Retrieval Baseline (BM25S)	39
3.2.3	Baseline Comparison Rationale	39
3.3	Error Analysis Methodology	40
3.3.1	Quantitative Analysis Framework	40
3.3.2	Query Length Bucketing	41
3.3.3	Failed Query Inspection	41
3.4	Query2Doc Implementation	42
3.4.1	Query2Doc Technique	42
3.4.2	Modifications from the Original Paper	43
3.4.3	LLM Configuration	44
3.4.4	Batch Processing Optimisation	44
3.4.5	Dense Retrieval with Query2Doc	45
3.4.6	BM25 Retrieval with Query2Doc	45
3.5	Model Comparison Methodology	46
3.5.1	Model Selection Criteria	46

3.5.2	Standardised Evaluation Protocol	47
3.5.3	Temperature Selection	48
3.5.4	Model-Specific Technical Issues	48
3.5.5	Quantisation Strategy	49
3.6	Query Repetition for Sparse Retrieval	49
3.7	Hybrid Retrieval Fusion	51
3.8	Corpus-Steered Query Expansion (CSQE)	52
3.8.1	The CSQE Pipeline	52
3.8.2	Component Ablation Design	54
3.8.3	Retriever-Specific Application	55
3.9	Per-Query Error Analysis	56
4	Results and Discussion	58
4.1	Baseline Results and Comparison	58
4.1.1	Complementary Strengths	59
4.2	Error Analysis Findings	60
4.2.1	Overall Failure Rate	60
4.2.2	Short Query Performance Gap	61
4.2.3	Retrieval Coverage Analysis	62
4.2.4	Technique Selection Rationale	63
4.3	Query2Doc Results	64
4.3.1	Dense Retrieval	64
4.3.2	BM25 Retrieval	65
4.3.3	Comparison with the Original Query2Doc Paper	66
4.4	Model Comparison Results	66
4.4.1	Dense Retrieval Leaderboard	67
4.4.2	BM25 Retrieval Results	69
4.4.3	Dropped Models Analysis	69
4.5	Key Findings and Analysis	70
4.5.1	Model Size Correlates with Dense Retrieval Performance	70
4.5.2	Generational Improvement Within Model Families	71
4.5.3	Arabic Specialisation vs. Multilingual Training	72
4.5.4	Dense vs. BM25 Behaviour with Query Enhancement	73

4.5.5	Model Selection Summary	73
4.6	BM25 Query Repetition Results	74
4.7	Hybrid Retrieval Fusion Results	77
4.8	Corpus-Steered Query Expansion Results	79
4.8.1	Main CSQE Results	79
4.8.2	Component Ablation	79
4.9	CSQE with Hybrid Fusion	80
4.9.1	Fusion Configuration Comparison	81
4.9.2	Overall System Progression	84
4.10	Per-Query Error Analysis	84
4.10.1	Win/Loss Distribution	85
4.10.2	First-Pass Quality as the Largest Modulator	86
4.10.3	Big Wins: The Corpus Grounding Effect	88
4.10.4	Regression Analysis	89
4.11	Summary of All Experiments	91
5	Conclusion and Recommendations	93
5.1	Conclusions	93
5.2	Challenges	96
5.3	Recommendations for Future Work	98
	Bibliography	100

List of Abbreviations

TE Test Example

AE Another Example

Chapter 1

Introduction

The rapid growth of digital Arabic content has intensified the need for intelligent information retrieval systems capable of accurately locating relevant information within large corpora. Retrieval-Augmented Generation (RAG) has emerged as the dominant paradigm for grounding Large Language Model (LLM) outputs in external knowledge, mitigating hallucination by retrieving relevant passages before generation. However, the quality of a RAG system’s final output is fundamentally bounded by the quality of its retrieval component: if the retriever fails to surface relevant documents, the generator cannot compensate for this deficiency regardless of its capabilities.

Arabic information retrieval faces challenges that compound this retrieval bottleneck. The Arabic language employs a root-and-pattern morphological system in which a single root can yield dozens of surface forms, creating severe vocabulary mismatch between queries and documents. Orthographic variations in characters such as Hamza and Alif introduce further inconsistencies, while the diglossic nature of Arabic—where Modern Standard Arabic (MSA) is used in formal text but regional dialects dominate informal communication—creates a systematic gap between how users formulate queries and how knowledge bases are written. These linguistic characteristics amplify the retrieval failure rate beyond what is observed for English and other morphologically simpler languages.

Query enhancement—the modification of a user’s query before it is submitted to the retriever—offers a promising pre-retrieval intervention. Rather than modifying the retriever or reindexing the knowledge base, query enhancement operates as a modular layer that enriches the query with additional vocabulary, context, and semantic information. This approach is retriever-agnostic, requires no changes to the existing infrastructure, and can be deployed incrementally. Recent work has demonstrated that LLMs can generate effective query expansions by producing pseudo-documents that bridge the vocabulary

gap between queries and relevant passages.

However, existing query enhancement research is overwhelmingly English-centric and relies on proprietary models with 175 billion or more parameters. The foundational techniques—Query2Doc, HyDE, and GRF—were all developed and validated using GPT-3 or comparable commercial APIs. Whether these techniques transfer effectively to Arabic, and whether small open-source models running on consumer-grade hardware can serve as viable alternatives, has not been systematically investigated.

1.1 Problem Definition

The problem addressed in this thesis lies at the intersection of three gaps in the current state of Arabic information retrieval.

The retrieval gap. Arabic information retrieval systems exhibit high failure rates even on standard benchmarks. Dense and sparse retrieval methods, while achieving reasonable aggregate performance, fail on a substantial proportion of individual queries. Short queries, which carry insufficient information for the retriever to identify relevant passages, are disproportionately affected. This information poverty in short queries is a fundamental limitation that cannot be addressed by improving the retriever alone.

The language gap. LLM-based query enhancement techniques were developed for English and validated on English benchmarks using English-language models. Arabic presents distinct challenges—morphological richness, orthographic variation, and vocabulary mismatch—that may cause these techniques to behave differently. The effectiveness of generative query enhancement for Arabic has not been empirically established.

The resource gap. The query enhancement literature relies almost exclusively on proprietary models with hundreds of billions of parameters, accessed through commercial APIs. This creates a practical barrier for deployment in resource-constrained settings. Whether small open-source models with 2–8 billion parameters, executable on freely available cloud GPUs, can generate useful Arabic query expansions remains an open

question.

These three gaps motivate the central research question of this thesis: *To what extent can small, open-source LLMs improve Arabic information retrieval through query enhancement, and what model characteristics determine effectiveness?*

1.2 Objectives

The objectives of this thesis are as follows:

1. To establish retrieval baselines for Arabic information retrieval using both dense and sparse methods, tested independently to isolate the sources of improvement from subsequent query enhancement experiments.
2. To conduct systematic error analysis on the baseline results to identify the primary failure patterns in Arabic retrieval and guide the selection of query enhancement techniques.
3. To adapt and implement the Query2Doc query enhancement technique for Arabic zero-shot application using small open-source LLMs, including engineering optimisations for practical deployment on consumer-grade GPUs.
4. To evaluate and compare ten open-source LLMs with 2–8 billion parameters for Arabic query expansion quality using a standardised evaluation protocol across both dense and sparse retrieval paradigms.
5. To analyse the interaction between LLM-based query enhancement and retrieval paradigm to identify model characteristics—such as parameter count, training data composition, and vocabulary design—that predict query expansion effectiveness for Arabic.
6. To investigate query repetition as a technique for resolving the BM25 term-dilution degradation observed when pseudo-documents are concatenated with short Arabic queries, and to identify the optimal repetition factor across multiple LLM models.

7. To adapt and evaluate Corpus-Steered Query Expansion (CSQE) for Arabic retrieval, grounding LLM-generated expansions in first-pass retrieved documents and assessing the contribution of corpus-grounded versus blind expansion components through ablation.
8. To combine corpus-steered query expansion with hybrid BM25–Dense retrieval and determine the optimal retriever-specific query representation strategy, examining whether applying query expansion asymmetrically to one retriever outperforms applying it to both.

1.3 Thesis Layout

The remainder of this thesis is organised as follows:

Chapter 2 establishes the theoretical foundations required to understand the research. It introduces the core concepts of Large Language Models, Retrieval-Augmented Generation, information retrieval methods (sparse, dense, and hybrid), and query enhancement techniques. The mathematical formulations underlying the retrieval and evaluation methods are presented, followed by detailed descriptions of all models used in the experimental work. The chapter concludes with a review of related literature and the identification of the research gap addressed by this thesis.

Chapter 3 presents the methodology employed throughout the research. It describes the dataset (MIRACL Arabic), hardware and software environment, and evaluation metrics. The baseline implementation for both dense (mDPR) and sparse (BM25S) retrieval is detailed, followed by the error analysis methodology that guided technique selection. The adaptation of Query2Doc for Arabic zero-shot application is described, including modifications from the original paper and engineering optimisations. The chapter then presents the query repetition methodology for resolving BM25 term-dilution, the hybrid BM25–Dense fusion methodology, and the Corpus-Steered Query Expansion (CSQE) pipeline including its component ablation design and retriever-specific application strategy. The chapter concludes with the per-query error analysis methodology.

Chapter 4 reports the results and discussion in zigzag correspondence with the methodology. Baseline performance is analysed, error analysis findings are presented, and Query2Doc results are reported for both dense and sparse retrieval. The comprehensive model comparison leaderboards are presented and analysed, including the examination of dropped models. The chapter then reports the expanded experimental results: query repetition sweep across nine models, hybrid retrieval fusion configurations, CSQE component ablation and alpha sweep, CSQE combined with hybrid fusion across three retriever assignment configurations, the overall system progression table, and a per-query error analysis decomposing the sources of improvement and regression.

Chapter 5 presents the overall conclusions, discusses the challenges encountered during the research, and provides recommendations for future work in Arabic query enhancement for RAG systems.

Chapter 2

Theoretical Background and Literature Review

This chapter establishes the theoretical foundations and reviews the current state-of-the-art relevant to the research presented in this thesis. Section 2.1 introduces the core concepts of Large Language Models (LLMs), Retrieval-Augmented Generation (RAG) systems, information retrieval methods, query enhancement techniques, and the specific challenges posed by the Arabic language. Section 2.2 presents the mathematical formulations underlying the retrieval and evaluation methods employed. Section 2.3 surveys candidate Arabic retrieval benchmarks and motivates the selection of the MIRACL Arabic dataset. Section 2.4 provides detailed descriptions of all models utilized in the experimental work. Finally, Section 2.5 surveys the related literature and identifies the research gap addressed by this thesis.

2.1 Theoretical Background

2.1.1 Large Language Models and the Transformer Architecture

Large Language Models (LLMs) are deep neural networks trained on vast corpora of text data to model the statistical properties of natural language. The modern era of LLMs was initiated by the introduction of the Transformer architecture by Vaswani et al. [1], which replaced recurrent processing with a self-attention mechanism that computes relationships between all positions in a sequence simultaneously.

The Transformer architecture consists of two main components: an encoder that

processes input sequences into contextual representations, and a decoder that generates output sequences autoregressively. The core innovation is the *scaled dot-product attention* mechanism, which computes attention weights between query (Q), key (K), and value (V) matrices as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V \quad (2.1)$$

where d_k is the dimensionality of the key vectors. Multiple attention heads are employed in parallel, allowing the model to attend to information from different representation subspaces.

Two dominant architectural paradigms have emerged from the original Transformer design. *Encoder-only* models such as Bidirectional Encoder Representations from Transformers (BERT) process entire sequences bidirectionally and are well-suited for classification and retrieval tasks. *Decoder-only* models, which underpin most modern LLMs, generate text autoregressively by predicting each token conditioned on all preceding tokens.

The development of LLMs typically involves two stages. During *pre-training*, the model is trained on trillions of tokens from diverse text sources using self-supervised objectives such as next-token prediction. This stage imbues the model with broad linguistic and world knowledge. During *instruction tuning* (also called Supervised Fine-Tuning, or SFT), the pre-trained model is further trained on curated prompt-response pairs to follow user instructions. Alignment techniques such as Direct Preference Optimization (DPO) and Reinforcement Learning from Human Feedback (RLHF) are subsequently applied to improve helpfulness and safety [2].

A fundamental limitation of LLMs is that their knowledge is *parametric*—encoded statically in model weights during training. This makes the knowledge prone to becoming outdated, incomplete for specialized domains, and susceptible to *hallucination*: the generation of fluent but factually incorrect text. The model cannot distinguish between information it has reliably learned and information it is confabulating, nor can its knowledge be updated without costly retraining [3].

2.1.2 Retrieval-Augmented Generation

Retrieval-Augmented Generation (RAG) was introduced by Lewis et al. [3] to address the limitations of purely parametric language models. RAG augments an LLM with access to an external knowledge base by incorporating a retrieval step before generation, thereby grounding the model’s outputs in verifiable sources.

The standard RAG architecture consists of two primary components:

- **The Retriever:** Given a user query q , the retriever searches a knowledge base $\mathcal{D} = \{d_1, d_2, \dots, d_N\}$ to identify a set of relevant passages $\{d_1, d_2, \dots, d_k\}$, ranked by relevance to q . Retrieval methods are discussed in detail in Section 2.1.3.
- **The Generator:** An LLM that produces an answer a conditioned on both the original query q and the retrieved passages. The retrieved context provides factual grounding, reducing the likelihood of hallucination and enabling the model to cite sources.

The RAG pipeline operates sequentially: (1) the user submits a query q ; (2) the retriever identifies the top- k relevant passages from the knowledge base; (3) the generator produces an answer conditioned on the query and retrieved passages; and (4) the system returns the generated answer, optionally with source citations. This architecture effectively separates knowledge storage (in the external corpus) from language generation (in the LLM), allowing the knowledge base to be updated independently without retraining the model.

The quality of the final generated answer is fundamentally bounded by the quality of the retrieved passages. If the retriever fails to surface relevant documents, the generator cannot compensate for this deficiency regardless of its capabilities. This *retrieval bottleneck* motivates the investigation of techniques to improve retrieval effectiveness, which is the central focus of this thesis. Figure 2.1 illustrates the architecture of the standard RAG framework described above.




Figure 2.1: High-level architecture of a Retrieval-Augmented Generation (RAG) system. A user query is sent to a retriever that consults the corpus and returns the top- k passages, which are passed alongside the query to an LLM to produce a context-augmented answer.

2.1.3 Information Retrieval Methods

Information retrieval methods can be broadly categorized into sparse, dense, and hybrid approaches, each with distinct strengths and limitations.

2.1.3.1 Sparse Retrieval

Sparse retrieval methods represent queries and documents as high-dimensional sparse vectors, where each dimension corresponds to a term in the vocabulary and the value reflects the term's importance. The most widely used sparse retrieval method is Best Matching 25 (BM25), a probabilistic ranking function based on the bag-of-words model [4]. BM25 computes relevance scores using term frequency, inverse document frequency, and document length normalization. The mathematical formulation of BM25 is presented in Section 2.2.1.

Sparse methods excel at exact term matching and are computationally efficient, requiring no neural network inference. However, they suffer from the *vocabulary mismatch problem*: if the query and document use different terms to express the same concept, sparse methods fail to identify the relevance. This limitation is particularly acute for morphologically rich languages such as Arabic, where a single concept may be expressed through numerous surface forms.

2.1.3.2 Dense Retrieval

Dense retrieval methods address the vocabulary mismatch problem by encoding queries and documents into continuous vector representations (embeddings) within a shared semantic space. Karpukhin et al. [5] introduced Dense Passage Retrieval (DPR), which employs a bi-encoder architecture: two independent BERT-based encoders produce fixed-size dense vectors for queries and passages respectively. Relevance is computed via vector similarity (typically dot product or cosine similarity, as formulated in Section 2.2.2) between the query embedding and passage embeddings.

Dense retrievers are trained on relevance-annotated data using contrastive learning objectives, where the model learns to place relevant query-passage pairs closer together in the embedding space while pushing irrelevant pairs apart. At inference time, approximate nearest neighbor search algorithms (e.g., FAISS) enable efficient retrieval over millions of passages.

While dense retrieval captures semantic relationships that sparse methods miss, it can struggle with rare terms, entity names, and highly specific queries where exact lexical matching is important [6].

2.1.3.3 Hybrid Retrieval

Hybrid retrieval combines sparse and dense methods to leverage their complementary strengths. Sparse retrieval excels at exact matching and handles rare terms well, while dense retrieval captures semantic similarity and handles paraphrasing. Two fusion strategies are employed in this thesis: Reciprocal Rank Fusion (RRF) [7], which combines ranked lists by summing inverse rank scores and requires no score normalisation, and score-level interpolation (Convex Combination, CC), which computes a weighted sum of min-max normalised sparse and dense scores. The formal definitions of both fusion functions are given with the other scoring functions in Section 2.2.3.

2.1.4 Query Enhancement Techniques

Query enhancement refers to a family of techniques that modify the user’s query before it is submitted to the retriever, with the goal of improving retrieval effectiveness. Song and Zheng [2] categorize query optimization methods into four atomic operations: *expansion*, *decomposition*, *disambiguation*, and *abstraction*. Each operation is described below; this thesis focuses on query expansion, the family most directly suited to the vocabulary-mismatch and information-poverty problems that dominate Arabic retrieval.

Query Expansion broadens the scope of a query by adding related terms, context, or generated text. Traditional query expansion methods such as Relevance Model 3 (RM3) extract expansion terms from pseudo-relevant documents retrieved in a first pass. Modern approaches leverage LLMs to generate expansion content, which can take several forms:

- **Hypothetical Document Embeddings (HyDE):** Proposed by Gao et al. [8], HyDE instructs an LLM to generate a hypothetical document that would answer the query. This generated document is then encoded using an unsupervised contrastive encoder, and the resulting embedding is used for retrieval. The encoder acts as a lossy compressor that filters out hallucinated details while retaining relevant semantic features.
- **Query2Doc:** Introduced by Wang et al. [6], this method generates a pseudo-document using an LLM and concatenates it with the original query to form an enriched query representation. For sparse retrieval, the original query is repeated multiple times before concatenation to prevent term dilution by the longer pseudo-document.
- **Generative Relevance Feedback (GRF):** Proposed by Mackie et al. [9], GRF uses an LLM to generate diverse text (keywords, essays, facts) relevant to the query. The generated text is treated as a feedback set from which expansion terms are extracted using a probabilistic model.

Query Decomposition breaks a complex or multi-hop query into simpler sub-queries that can be retrieved separately and recombined. Chan et al. [33] train a language model

(RQ-RAG) to decompose, rewrite, and disambiguate queries on demand, improving retrieval for multi-hop question answering where a single retrieval pass is insufficient.

Query Disambiguation reformulates an ambiguous or poorly phrased query into a clearer, more retrievable form without necessarily adding new terms. The Rewrite-Retrieve-Read framework [10] employs a trainable language model to rewrite queries before retrieval, optimising the rewriter using reinforcement-learning signals from the downstream reader’s performance.

Query Abstraction steps back from a specific query to a more general question whose answer is easier to retrieve and which then guides retrieval for the original query. Zheng et al. [?] introduced *step-back prompting*, in which the model first derives a higher-level concept or principle before answering, improving retrieval and reasoning on knowledge-intensive tasks.

Among these four families, query expansion is adopted in this thesis because it directly addresses the dominant Arabic retrieval failure modes (vocabulary mismatch and short, information-poor queries) while remaining model-agnostic and requiring no task-specific training.

These query enhancement techniques operate as modular layers that can be added to existing RAG systems without modifying the retriever index or the generator, making them particularly attractive for practical deployment. Table 2.1 summarises the query-enhancement literature reviewed for this thesis, organised by method family.

Table 2.1: Query-enhancement papers reviewed for this thesis, with method family, evaluation dataset, and headline result.

Paper	Year	Family	Method	Dataset
Query2Doc (Wang et al.)	2023	generation	LLM pseudo-document concatenation	MS MARCO + BEIR
HyDE (Gao et al.)	2022	generation	LLM hypothetical document	TREC-DL / BEIR
GAR (Mao et al.)	2021	generation	LLM generates answer/context for query rewriting	open-domain QA
GRF (Mackie et al.)	2023	generation	Generative Relevance Feedback for dense	TREC-DL
CSQE (Lei et al.)	2024	generation	Corpus-Steered Query2Doc	TREC-DL
MuGI (Zhang et al.)	2024	generation	Multi-Text Generation Integration	BEIR
ThinkQE (Lei et al.)	2025	generation	Iterative chain-of-thought expansion	TREC-DL
Exp4Fuse (Liu et al.)	2025	fusion	Original + expanded query via modified RRF	multiple
DAPR (Wang et al.)	2024	generation	Dense Approach via Pre-trained Retrievers	academic IR
AQE (Yang et al.)	2025	generation	Aligned Query Expansion via LLM alignment	academic IR
KAR (Xia et al.)	2025	generation	Knowledge-Aware Query Expansion	textual + relational retrieval
Song & Zheng (survey)	2024	survey	Query optimisation taxonomy (4 ops)	–
QE survey (arXiv 2509.07794)	2025	survey	Comprehensive QE survey in LLM era	–

2.1.5 Arabic Language Processing Challenges

The Arabic language presents several characteristics that compound the retrieval challenges described in the preceding sections.

Morphological Richness. Arabic employs a root-and-pattern morphological system in which words are derived from typically tri-consonantal roots through templatic patterns and extensive affixation. A single root can yield dozens of surface forms through prefixes, suffixes, infixes, and clitics. For example, the root *k-t-b* (- - , “to write”) generates forms such as *kitaab* (book), *kaatib* (writer), *maktaba* (library), and *maktuub* (written). This morphological productivity creates a severe vocabulary mismatch problem, as queries and relevant documents may use different inflected forms of the same root.

Diglossia. Arabic exhibits a diglossic situation in which Modern Standard Arabic (MSA) is used in formal writing, education, and media, while various regional dialects (Egyptian, Levantine, Gulf, Maghrebi, and others) are used in daily speech and informal written communication. Knowledge bases are predominantly composed of MSA text, whereas user queries—particularly in conversational settings—may be formulated in dialectal Arabic, creating a systematic gap between query and document vocabularies [11].

Orthographic Variations. Arabic script permits multiple orthographic representations for certain characters, leading to inconsistencies even within MSA. Common variations include the different forms of Hamza (, (, (, and Ya/Alif Maqsura (. (. These variations can cause both sparse and dense retrieval methods to treat orthographic variants as distinct terms.

Diacritical Marks. Arabic diacritics (*tashkeel*) disambiguate pronunciation and meaning but are typically omitted in modern texts. This omission increases lexical ambiguity, as identical consonantal skeletons may correspond to multiple words with different meanings.

These characteristics collectively create what may be termed a “morphological gap” between user queries and knowledge base documents, which is distinct from and often more severe than the vocabulary mismatch observed in English information retrieval.

2.2 Mathematical Models

This section presents the mathematical formulations of the retrieval scoring functions and evaluation metrics employed in this thesis.

2.2.1 BM25 Scoring Function

The BM25 ranking function [4] computes the relevance score of a document d with respect to a query $q = \{q_1, q_2, \dots, q_n\}$ consisting of n terms:

$$\text{BM25}(q, d) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{f(q_i, d) \cdot (k_1 + 1)}{f(q_i, d) + k_1 \cdot \left(1 - b + b \cdot \frac{|d|}{\text{avgdl}}\right)} \quad (2.2)$$

where $f(q_i, d)$ is the term frequency of query term q_i in document d , $|d|$ is the length of document d in tokens, avgdl is the average document length across the collection, k_1 controls the term frequency saturation (typically $k_1 = 1.2$), and b controls the degree of

document length normalization (typically $b = 0.75$).

The Inverse Document Frequency (IDF) component is computed as:

$$\text{IDF}(q_i) = \ln \left(\frac{N - n(q_i) + 0.5}{n(q_i) + 0.5} + 1 \right) \quad (2.3)$$

where N is the total number of documents in the collection and $n(q_i)$ is the number of documents containing term q_i . The IDF assigns higher weight to rarer terms, which are generally more discriminative for retrieval.

2.2.2 Dense Retrieval and Cosine Similarity

In dense retrieval, queries and documents are encoded into dense vector representations by neural encoder models. Given a query encoder E_q and a document encoder E_d , the relevance score between query q and document d is computed as the cosine similarity between their embedding vectors:

$$\text{sim}(q, d) = \cos(\mathbf{e}_q, \mathbf{e}_d) = \frac{\mathbf{e}_q \cdot \mathbf{e}_d}{\|\mathbf{e}_q\| \|\mathbf{e}_d\|} \quad (2.4)$$

where $\mathbf{e}_q = E_q(q)$ and $\mathbf{e}_d = E_d(d)$ are the embedding vectors for the query and document respectively, $\mathbf{e}_q \cdot \mathbf{e}_d$ denotes the dot product, and $\|\cdot\|$ denotes the Euclidean norm. When embedding vectors are L_2 -normalized, cosine similarity is equivalent to the dot product, and retrieval can be performed efficiently using Maximum Inner Product Search (MIPS) algorithms.

2.2.3 Hybrid Fusion Functions

Hybrid retrieval (Section 2.1.3.3) combines the ranked lists of a sparse and a dense retriever using one of two fusion functions. Reciprocal Rank Fusion (RRF) [7] combines ranked lists by summing inverse rank scores:

$$s_{\text{RRF}}(d) = \sum_{r \in \mathcal{R}} \frac{1}{k + \text{rank}_r(d)} \quad (2.5)$$

where $\mathcal{R}$ is the set of ranked lists being fused and k is a smoothing constant (typically $k = 20$) that reduces the impact of top-ranked documents. RRF is parameter-free in the sense that it does not require score normalisation, making it robust to score distribution differences between sparse and dense retrievers.

Score-level interpolation (Convex Combination, CC) computes a weighted sum of normalised retriever scores:

$$s_{\text{CC}}(d, q) = \alpha \cdot \hat{s}_{\text{Dense}}(d, q) + (1 - \alpha) \cdot \hat{s}_{\text{BM25}}(d, q) \quad (2.6)$$

where $\hat{s}$ denotes min-max normalised scores and $\alpha \in [0, 1]$ controls the relative contribution of the dense retriever. CC is sensitive to the normalisation strategy and score scale differences between the two systems.

2.2.4 Evaluation Metrics

Three standard information retrieval metrics are used throughout this thesis to evaluate retrieval performance.

2.2.4.1 Recall at k (Recall@ k)

Recall at k measures the proportion of all relevant documents that appear within the top- k retrieved results:

$$\text{Recall@}k = \frac{|\{\text{relevant documents}\} \cap \{\text{top-}k \text{ retrieved}\}|}{|\{\text{relevant documents}\}|} \quad (2.7)$$

Recall@ k captures the retriever's ability to surface relevant documents within the top- k positions. In RAG systems, Recall@ k at the retrieval stage directly bounds the information

available to the generator.

2.2.4.2 Normalized Discounted Cumulative Gain at k (NDCG@ k)

Normalized Discounted Cumulative Gain (NDCG) measures ranking quality by considering both the relevance of retrieved documents and their positions. Documents at higher ranks contribute more to the score, reflecting the intuition that users examine top-ranked results more carefully.

The Discounted Cumulative Gain at position k is defined as:

$$\text{DCG@}k = \sum_{i=1}^k \frac{2^{\text{rel}_i} - 1}{\log_2(i + 1)} \quad (2.8)$$

where rel_i is the relevance grade of the document at position i . The ideal DCG (IDCG@ k) is computed by sorting all relevant documents by decreasing relevance. NDCG@ k is the ratio:

$$\text{NDCG@}k = \frac{\text{DCG@}k}{\text{IDCG@}k} \quad (2.9)$$

NDCG@ k ranges from 0 to 1, where 1 indicates a perfect ranking.

2.2.4.3 Mean Reciprocal Rank (MRR)

Mean Reciprocal Rank measures the average of the reciprocal ranks of the first relevant document across a set of queries Q :

$$\text{MRR} = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \frac{1}{\text{rank}_i} \quad (2.10)$$

where rank_i is the position of the first relevant document for the i -th query. MRR emphasizes the ability to place at least one relevant document at the top of the ranked

list.

2.3 Evaluation Dataset Selection

The choice of evaluation benchmark is central to a query enhancement study, because query enhancement is only meaningful where the benchmark exposes a genuine query–document matching problem. Several Arabic and multilingual benchmarks were therefore surveyed and assessed against five criteria: (1) a *retrieval-first* design with explicit relevance judgments, rather than a reading-comprehension task in which retrieval is incidental; (2) *natural, native-authored* queries rather than machine-translated or template-generated ones, so that the query–document mismatch is linguistically authentic; (3) coverage of *Modern Standard Arabic* over an open-domain corpus; (4) sufficient *scale* for stable evaluation while remaining tractable on the available compute; and (5) *open availability* for reproducibility. Table 2.2 summarises the principal candidates.

Table 2.2: Arabic retrieval and question-answering datasets surveyed for benchmark selection. “Suitability” refers to fitness for studying query enhancement for retrieval.

Dataset	Ar. queries	Variety	Query origin	Primary challenge	Suitability
MIRACL [36]	~8,700	MSA	Native (ad-hoc)	Retrieval recall	High
ArabicaQA [?]	~89,000	MSA	Crowd (paraphrased)	Scale / difficulty	High
TyDi QA [?]	~23,000	MSA	Native (info-seeking)	Morphology / typology	Medium
MultiNativQA [?]	~12,000	MSA + dialect	Native (cultural)	Dialect / cultural gap	Medium
Mintaka [?]	~20,000	MSA (transl.)	Translated	Multi-hop reasoning	Medium
ARCD [?]	~1,400	MSA	Crowd	Small scale	Low

Selected benchmark: MIRACL (Arabic). MIRACL [36] was selected as the primary evaluation benchmark. Unlike reading-comprehension datasets, it is purpose-built for ad-hoc retrieval and provides human relevance judgments produced by native-speaker annotators, including *hard negatives*—passages that share surface terms with the query but are not relevant. This makes it a direct test of whether an enhanced query can separate the truly relevant passage from lexically similar distractors, which is precisely

the behaviour query enhancement aims to improve. Its queries are written by native speakers from genuine information needs, avoiding the priming and translation artefacts that affect translated datasets; its corpus is the standard open-domain Arabic Wikipedia; and its development set of 2,896 queries is large enough for stable evaluation yet tractable under the compute constraints of this work.

Alternatives considered. ArabicaQA [?] offers an order of magnitude more queries and an explicit “easy”/“difficult” split by retrieval rank, making it attractive for stress-testing enhancement on hard queries; TyDi QA [?], whose information-seeking annotation methodology MIRACL builds upon, contributes highly natural queries; and MultiNativQA [?] uniquely includes regional Arabic dialects, making it the natural vehicle for evaluating dialect-aware query rewriting. Mintaka [?] targets multi-hop reasoning suited to query decomposition but is professionally *translated* into Arabic, introducing translationese. Machine-translated (Arabic-SQuAD) and template-generated (ACQAD) resources were excluded outright, as their query–document mismatch is an artefact of translation or generation rather than authentic Arabic phrasing, which would reward enhancement methods for the wrong reasons.

Limitation. MIRACL covers only Modern Standard Arabic over Wikipedia and therefore does not exercise dialectal query enhancement. This scoping choice is revisited as a direction for future work in Chapter 5, where dialect-rich resources such as MultiNativQA are proposed for evaluating dialect-aware enhancement.

2.4 Models Used in Experiments

This section describes all language models employed in the experimental work of this thesis. The models were selected based on three primary criteria: (1) strong Arabic language support, either through specialized Arabic training or robust multilingual coverage; (2) compatibility with consumer-grade GPU hardware, specifically the NVIDIA T4 (15 GB VRAM) available through Google Colab and the NVIDIA A100 (40 GB VRAM); and (3) open-source availability with permissive licensing to ensure reproducibility [2]. Models are

organized into Arabic-specialized, multilingual, and experimental categories. Table 2.3 summarises the ten LLMs evaluated; the remainder of this section describes each in detail.

Table 2.3: Open-source LLMs evaluated for Arabic query enhancement in this thesis, ordered by parameter count.

Model	Family	Params (B)	Multilingual	Arabic-specialised	Quantisation
SILMA Kashif-2B	Silma AI	2	No	Yes (RAG)	FP16
Qwen 2.5 3B	Alibaba	3	Yes (29 langs)	No	FP16
Falcon-H1-3B	TII	3	Yes	Yes	BF16
Qwen3-4B	Alibaba	4	Yes (119 langs)	No	FP16
Gemma 3 4B	Google	4	Yes	No	FP16/BF16
Qwen 2.5-7B	Alibaba	7	Yes (29 langs)	No	FP16
ALLaM-7B (DROPPED)	NCAI	7	Yes	Yes	FP16
Jais-2-8B	G42	8	Yes	Yes	BF16
Qwen3-8B	Alibaba	8	Yes (119 langs)	No	FP16
Aya Expanse 8B	Cohere	8	Yes (101 langs)	No	4-bit NF4

2.4.1 Arabic-Specialized Models

2.4.1.1 Falcon-H1-Arabic-3B

Falcon-H1-Arabic-3B is a 3.15 billion parameter language model developed by the Technology Innovation Institute (TII), Abu Dhabi [13]. Unlike conventional Transformer-based LLMs, Falcon-H1 employs a hybrid Mamba2-Transformer architecture that combines State Space Model (SSM) layers [14] with standard multi-head attention layers and SwiGLU activation functions. The model contains 32 layers with a hidden size of 2,560, utilizing 32 Mamba heads alongside 10 attention heads with Grouped Query Attention (GQA, 2 key-value heads).

The hybrid architecture was designed to combine the linear-time sequence processing efficiency of SSMs with the strong in-context learning capabilities of attention mechanisms. The model supports a context length of 128K tokens and was purpose-built for Arabic, listing it as a primary language alongside 17 additional languages.

On the Open Arabic LLM Leaderboard (OALL v2), Falcon-H1-Arabic-3B achieved an average score of approximately 62%, roughly 10 points ahead of competing models at the same parameter count. The model was released in May 2025 under an open license. In FP16 precision, the model requires approximately 10–11 GB of VRAM including SSM state buffers, fitting on both T4 and A100 GPUs without quantization.

2.4.1.2 Jais-2-8B

Jais-2-8B-Chat is an 8.09 billion parameter Arabic-centric language model developed collaboratively by Inception (G42), Mohamed bin Zayed University of Artificial Intelligence (MBZUAI), and Cerebras [15]. The model employs a standard decoder-only Transformer architecture with Rotary Position Embeddings (RoPE), Squared-ReLU activation, and LayerNorm. It contains 32 layers with a hidden size of 3,328 and 26 attention heads.

A distinguishing feature of Jais-2 is its Arabic-centric vocabulary of 150,272 tokens, trained from scratch to optimize Arabic tokenization efficiency. The model was pre-trained on 2.6 trillion tokens—6.6 times more data than its predecessor Jais-1—with substantial Arabic and English representation. The instruction-tuned variant was refined through a three-stage alignment pipeline comprising large-scale Supervised Fine-Tuning (SFT) on Arabic and English prompt-response pairs, followed by DPO, and finally Group Relative Policy Optimization (GRPO).

On the AraGen-12-24 benchmark, Jais-2-8B-Chat achieved an overall 3C3H score of 58.64%, outperforming Fanar-1-9B and ALLaM-7B. The model was released in December 2024 under Apache 2.0 license (gated access). It requires BF16 precision due to Squared-ReLU activation overflow in FP16, using approximately 16.6 GB of VRAM on the A100.

2.4.1.3 ALLaM-7B

ALLaM-7B-Instruct-preview is a 7 billion parameter Arabic-English language model developed by the National Center for Artificial Intelligence (NCAI) at the Saudi Data and Artificial Intelligence Authority (SDAIA) [16]. Published at the International Conference on Learning Representations (ICLR) 2025, ALLaM employs the LlamaForCausalLM

architecture with an expanded vocabulary of 64,000 tokens (doubled from Llama-2’s 32K to accommodate Arabic).

The model was trained from scratch in two phases: (1) a foundation phase on 4 trillion English tokens, followed by (2) continued pre-training on 1.2 trillion mixed Arabic-English tokens at a 45%/45%/10% Arabic/English/code ratio, totaling 5.2 trillion training tokens. A substantial share of the Arabic pre-training data was derived from machine-translated English sources. Instruction tuning and DPO-based alignment were subsequently applied.

On the Arabic Massive Multitask Language Understanding (MMLU) benchmark, ALLaM achieved a 0-shot score of 67.78. The model was released in February 2025 under Apache 2.0 license. Notably, the publicly available version carries “preview” (alpha) status, indicating that it may not fully reflect the capabilities described in the published paper. The model requires approximately 14.5 GB of VRAM in FP16 on the A100.

2.4.1.4 SILMA Kashif-2B

SILMA Kashif-2B-Instruct is a 2 billion parameter language model developed by SILMA AI, an Arabic AI startup, with a specific focus on Arabic RAG applications [17]. The model was optimized for extractive question answering in Arabic, achieving a RAG Question Answering (RAGQA) benchmark score of 0.3478, which placed it at the top of the 3–9 billion parameter range on this benchmark.

The model supports a context length of 12K tokens and operates in FP16 precision, requiring approximately 4–5 GB of VRAM. Its small size makes it the most computationally efficient model in the comparison, suitable for deployment on consumer-grade hardware without quantization.

2.4.2 Multilingual Models

2.4.2.1 Qwen 2.5 3B

Qwen 2.5-3B-Instruct is a 3.09 billion parameter multilingual language model developed by the Qwen Team at Alibaba Cloud [18]. The model employs a standard decoder-only Transformer architecture with GQA (16 query heads, 2 key-value heads), RMSNorm, SiLU activation, and a vocabulary of 151,936 tokens. It was pre-trained on approximately 18 trillion tokens spanning 29 or more languages, with Arabic included as a supported language.

Qwen 2.5-3B-Instruct was released in September 2024 under the Qwen Research License (non-commercial use). The model requires approximately 6 GB of VRAM in FP16, fitting comfortably on the T4 GPU. In the context of this thesis, Qwen 2.5-3B served as the reference model for the first Query2Doc experiment, establishing the initial query enhancement baseline.

2.4.2.2 Qwen 2.5 7B

Qwen 2.5-7B-Instruct is the 7 billion parameter variant in the same model family [18]. It shares the same architecture and tokenizer as the 3B variant but with a wider hidden dimension and more attention heads, resulting in stronger performance across multilingual benchmarks (MMLU: 74.2). The model requires 4-bit quantization to fit on the T4 GPU but operates in full FP16 precision on the A100. Pre-quantized versions are available from both the Qwen team and the Unsloth community.

2.4.2.3 Qwen3-4B

Qwen3-4B is a 4.0 billion parameter (3.6 billion non-embedding) next-generation language model from the Qwen Team [19]. Released in May 2025, it represents a significant architectural evolution over Qwen 2.5: the number of supported languages and dialects was expanded from 29 to 119 (with broader Arabic coverage), training data was doubled

to approximately 36 trillion tokens, and attention capacity was increased to 32 query heads with 8 key-value heads.

A notable feature of Qwen3 is its built-in *thinking mode*, which generates internal reasoning within special tags before producing the final answer. For query expansion tasks, this thinking mode must be explicitly disabled to prevent reasoning artifacts from contaminating the generated pseudo-documents. The model also requires sampling-based decoding; greedy decoding causes performance degradation and repetitive output.

Qwen3-4B matches or exceeds Qwen 2.5-7B on many benchmarks despite having fewer parameters, achieving an MMMLU (multilingual MMLU) score of 71.42. It operates in FP16 precision requiring approximately 8.5 GB of VRAM.

2.4.2.4 Qwen3-8B

Qwen3-8B is the 8 billion parameter variant of the Qwen3 family [19], sharing the same architectural innovations as Qwen3-4B including 119-language support, thinking mode, and 36 trillion training tokens. As the larger variant, it provides stronger language understanding and generation capabilities. Like Qwen3-4B, thinking tags must be stripped from generated output when used for query expansion. The model operates in FP16 on the A100 GPU.

2.4.2.5 Gemma 3 4B-IT

Gemma 3 4B-IT is a 4 billion parameter instruction-tuned language model developed by Google DeepMind [20]. Released in March 2025, it supports over 140 languages and features a context length of 128K tokens. The model employs a standard dense Transformer architecture with multimodal capabilities (vision and text), though only the text modality was used in this thesis.

Despite its broad language coverage, Gemma 3 exhibited the weakest Arabic performance among the tested models, scoring approximately 10 points below Falcon-H1-Arabic-3B on the OALL benchmark. It serves as a lower-bound comparison point in the model

evaluation, representing a general-purpose multilingual model without Arabic-specific optimization.

2.4.2.6 Aya Expanse 8B

Aya Expanse 8B is an 8 billion parameter multilingual language model developed by Cohere Labs (CohereForAI), covering 23 languages with Arabic among the explicitly supported languages [21]. The model was trained using a data-centric approach that emphasizes high-quality multilingual training data, and its developers report that it outperforms comparable open-weight models, including Llama-3.1-8B-Instruct, on multilingual evaluations.

Aya Expanse employs a standard Transformer architecture and was loaded with 4-bit NF4 quantization in the experiments due to its 8 billion parameter size. Despite the quantization, it demonstrated strong Arabic generation quality, suggesting that architecture and training data quality outweigh precision for this task. The model is distributed under a CC-BY-NC-4.0 (non-commercial) license.

2.4.3 Retrieval Models

2.4.3.1 Multilingual Dense Passage Retriever (mDPR)

The Multilingual Dense Passage Retriever (mDPR) [5] is a bi-encoder dense retrieval model based on multilingual BERT. In this thesis, the pre-built mDPR index for MIRACL Arabic was used, which encodes the 2.1 million Arabic Wikipedia passages into dense vectors. mDPR was selected as the dense baseline because it was *not* fine-tuned on MIRACL training data, unlike models such as BGE-M3 and E5 which were partially trained on MIRACL. This choice provides a more realistic evaluation of query enhancement effectiveness, as improvements cannot be attributed to the retriever having memorized the evaluation queries.

2.4.3.2 BM25S

BM25S [23] is a Python-native implementation of the BM25 algorithm that eliminates the Java dependencies required by traditional implementations such as Pyserini. In this thesis, BM25S was configured with $k_1 = 0.9$ and $b = 0.4$, matching the Pyserini default configuration used in the official MIRACL benchmark to ensure comparability. (The BM25S library defaults are $k_1 = 1.5$, $b = 0.75$.) BM25S was selected over Pyserini for its simpler deployment pipeline and faster iteration speed, achieving 96% of the Pyserini BM25 baseline performance on the MIRACL Arabic evaluation set.

Notation. Throughout the remainder of this thesis, the term “BM25” refers to the BM25S implementation with the configuration above ($k_1 = 0.9$, $b = 0.4$), unless explicitly stated otherwise.

2.4.4 Model Comparison Summary

Table 2.4 presents a comparative summary of all language models used in the experiments. The models span a range of parameter counts from 2 billion to 8 billion, encompassing Arabic-specialized and multilingual architectures.

Table 2.4: Summary of language models used in the experiments.

Model	Developer	Params	Architecture	Arabic Focus	License
<i>Arabic-Specialized</i>					
Falcon-H1-3B	TII	3.15B	Hybrid Mamba2-Trans.	Primary	Open
Jais-2-8B	MBZUAI/Inception	8.09B	Decoder-only Trans.	Primary	Apache 2.0
ALLaM-7B	SDAIA	7.0B	LlamaForCausalLM	Primary	Apache 2.0
SILMA Kashif-2B	SILMA AI	2.0B	Decoder-only Trans.	Primary	Open
<i>Multilingual</i>					
Qwen 2.5 3B	Alibaba	3.09B	Decoder-only Trans.	29+ langs	Qwen Resea
Qwen 2.5 7B	Alibaba	7.0B	Decoder-only Trans.	29+ langs	Apache 2.0
Qwen3-4B	Alibaba	4.0B	Decoder-only Trans.	119 langs	Apache 2.0
Qwen3-8B	Alibaba	8.0B	Decoder-only Trans.	119 langs	Apache 2.0
Gemma 3 4B	Google DeepMind	4.0B	Decoder-only Trans.	140+ langs	Gemma TO
Aya Expanse 8B	Cohere Labs	8.0B	Decoder-only Trans.	101 langs	CC-BY-NC

2.5 Related Work

This section reviews the published literature on query enhancement for information retrieval and RAG systems, organized chronologically and thematically. The review progresses from foundational work on generative query expansion to modern approaches using smaller LLMs, and concludes with Arabic-specific information retrieval research.

2.5.1 Foundational Query Enhancement Research (2022–2023)

The foundational work on LLM-based query enhancement was conducted using large proprietary models. Gao et al. [8] introduced HyDE, which demonstrated that an instruction-following LLM (InstructGPT, 175B parameters) could generate hypothetical documents for zero-shot dense retrieval, achieving performance comparable to supervised models without requiring any relevance labels. HyDE’s key insight was that document-to-document similarity is more reliable than query-to-document similarity, and that even hallucinated documents capture useful relevance patterns when encoded through a dense bottleneck.

Wang et al. [6] extended this concept with Query2Doc, which generates pseudo-documents and concatenates them with the original query rather than replacing it. Using `text-davinci-003` (175B parameters) with few-shot prompting, Query2Doc achieved 3–15% improvement over BM25 and made sparse retrieval competitive with state-of-the-art dense methods. Critically, the authors found smaller models insufficient for quality query expansion at the time—a finding that has since been challenged by newer, more capable small models. For sparse retrieval, the paper introduced query repetition (repeating the original query 5 times before concatenation) to prevent term dilution.

Mackie et al. [9] proposed Generative Relevance Feedback (GRF), which uses GPT-3 to generate diverse content types (keywords, essays, facts, news articles) as feedback for query expansion. GRF achieved 5–19% improvement in MAP and 17–24% in NDCG@10 over BM25+RM3 baselines. In subsequent work [24], the authors extended GRF to dense and learned sparse retrieval models, demonstrating approximately 9–10% improvement in NDCG@20. A key finding was that GRF and traditional pseudo-relevance feedback are

complementary: their fusion achieved the best Recall@1000 in 17 out of 18 experiments.

Ma et al. [10] introduced the Rewrite-Retrieve-Read framework, which employs a trainable T5-large model as a query rewriter optimized via reinforcement learning from the downstream reader’s performance. This work demonstrated that small, trainable models can effectively improve the performance of large, frozen LLMs by operating at the query formulation stage.

A common characteristic of all foundational work is the reliance on 175B-parameter models (the GPT-3 family) for query expansion, with smaller models either untested or found inadequate.

2.5.2 Modern LLM-Based Query Enhancement (2024–2025)

The period from 2024 to 2025 saw significant progress in demonstrating that moderately-sized open-source LLMs (7–8B parameters) can perform effective query expansion, overturning the assumption that 175B-scale models are necessary.

Lei et al. [25] proposed Corpus-Steered Query Expansion (CSQE), which uses the target corpus to guide LLM-generated expansions. The CSQE pipeline operates in two stages. First, a first-pass sparse retrieval step retrieves a small set of top- k documents from the target corpus using the original query. These retrieved documents are then provided to the LLM alongside the query, instructing it to extract and synthesise topically relevant vocabulary and context grounded in the actual corpus content. This corpus-grounded expansion is combined with a blind expansion (standard Query2Doc) to produce the final enriched query. The key motivation is that corpus grounding anchors the LLM’s output to attested Wikipedia vocabulary, preventing the hallucination of plausible but factually incorrect entities that characterises blind generation for niche or ambiguous queries. Lei et al. reported that even a 7B-parameter model achieves a 30% improvement in mAP over BM25 on English benchmarks [25].

Zhang et al. [26] studied best practices for LLM-based query expansion and proposed MuGI (Multi-Text Generation Integration), which prompts an LLM to generate multiple pseudo-documents and integrates them with the original query. To prevent the lengthy

generated text from overwhelming the original query terms in sparse retrieval, MuGI employs an adaptive query-repetition strategy that scales the number of query repetitions to the length of the generated text—the basis for the adaptive repetition parameter (β) adopted in this thesis (Section 3.6).

Several additional studies have examined query enhancement with smaller and more efficient models. Xia et al. [29] proposed knowledge-aware query expansion, augmenting LLM-generated expansions with structured knowledge to improve both textual and relational retrieval. Yang et al. [30] demonstrated aligned query expansion, using LLM alignment to make query expansions more effective and efficient for information retrieval. Lei et al. [31] proposed ThinkQE, which refines query expansions through an evolving, reasoning-driven generation process. Zhang et al. [27] explored personalized, user-centric query expansion, reporting retrieval improvements on a personalization benchmark. Not all findings have been uniformly positive, however: Yoon et al. [28] caution that a substantial portion of the apparent gains from LLM-based query expansion on common benchmarks may stem from benchmark contamination (“knowledge leakage”) rather than genuine retrieval improvement, underscoring the importance of careful evaluation.

Research on knowledge distillation further reduced the model size requirements. Young et al. [32] proposed GaQR, which distills query rewriting capabilities from GPT-3.5 into Llama2-7B via a three-phase pipeline (generation, verification, distillation), achieving substantially faster inference than chain-of-thought generation methods. Chan et al. [33] introduced RQ-RAG, which trains Llama2-7B to autonomously refine queries through rewriting, decomposition, and disambiguation, achieving state-of-the-art performance for 7B-parameter models on multi-hop QA tasks. Bhusal [34] further demonstrated that knowledge distillation from GPT-3.5 to Flan-T5-base (approximately 250M parameters) can produce effective query augmentation, suggesting a path toward extremely efficient deployment.

Zhang et al. [35] addressed query robustness with QE-RAG, a benchmark specifically designed to evaluate RAG system performance under query entry errors (typos, misspellings, visual similarities). Their findings demonstrated that query enhancement techniques significantly improve retrieval robustness against noisy inputs—a characteristic analogous to the morphological and orthographic variations present in Arabic queries.

Collectively, this body of work demonstrates that the model size requirement for effective query expansion has decreased dramatically: from 175B parameters in 2022–2023 to 7–8B parameters in 2024–2025, with further compression possible via knowledge distillation. However, a critical gap remains: *no prior study has evaluated sub-7B models for zero-shot query expansion in monolingual Arabic retrieval*. The closest concurrent work, by Macmillan-Scott et al. [?], applies sub-7B models (including Gemma 3 4B) to the distinct *cross-lingual* Arabic retrieval setting; notably, it independently corroborates several observations of this thesis, including that reasoning-style meta-text can pollute sparse retrieval and that zero-shot prompting is preferable for short queries.

2.5.3 Arabic Information Retrieval and RAG

Research on Arabic-specific RAG systems has begun to establish baselines, though work on query enhancement for Arabic remains scarce.

Zhang et al. [36] introduced MIRACL (Multilingual Information Retrieval Across a Continuum of Languages), a large-scale multilingual retrieval benchmark covering 18 languages including Arabic. The Arabic subset contains over 2.1 million Wikipedia passages with human-annotated relevance judgments for 2,896 development queries and 799 test queries, making it the primary evaluation resource for Arabic retrieval research.

Alsubhi et al. [12] conducted the first systematic analysis of RAG pipeline components for Arabic. Testing chunking strategies, embedding models, rerankers, and generators, they found that sentence-aware chunking outperformed fixed-size and semantic chunking for Arabic text, and that multilingual embedding models (BGE-M3, multilingual E5-large) outperformed Arabic-specific models for retrieval. They also identified Aya-8B as a superior generator for semantically dense Arabic content compared to StableLM-1.6B.

El-Beltagy and Abdallah [11] explored RAG applications for Arabic, benchmarking twelve embedding models and five LLMs. Their work revealed that E5 and BGE models exhibited remarkable robustness to Arabic dialectal variations in queries, and that open-source models (Llama 3, Mistral 7B) performed comparably to GPT-3.5 Turbo for Arabic text generation in RAG contexts.

These studies establish that Arabic RAG systems can achieve reasonable performance with existing multilingual tools. However, they focus primarily on the retrieval and generation components of the pipeline. The query formulation stage—specifically, whether LLM-based query enhancement techniques developed for English transfer effectively to Arabic—remains uninvestigated.

2.5.4 Research Gap

The literature review reveals a clear gap at the intersection of two well-studied areas:

- a. **LLM-based query enhancement** has been extensively studied for English, with recent work demonstrating that 7–8B parameter models are sufficient for effective query expansion (Section 2.5.2). However, none of these studies evaluate query expansion for *monolingual* Arabic retrieval; the sole concurrent exception [?] addresses the distinct *cross-lingual* setting.
- b. **Arabic RAG systems** have established retrieval and generation baselines (Section 2.5.3), but query enhancement techniques—which operate upstream of both retrieval and generation—have not been investigated for Arabic.

Specifically, the following questions remain unanswered:

- Can modern open-source LLMs (2–8B parameters) perform effective zero-shot query expansion for Arabic information retrieval?
- How do Arabic-specialized models compare against multilingual models for this task?
- Does the Query2Doc technique, originally validated with 175B-parameter models on English text, transfer to Arabic with smaller models?
- What model characteristics (parameter count, architecture, Arabic training data volume, vocabulary design) most influence query expansion quality for Arabic?

A further gap exists in the application of corpus-steered expansion to non-English retrieval. The original CSQE work [25] was evaluated exclusively on English benchmarks using English-language models; whether corpus grounding provides the same benefits for Arabic—where BM25 homonym sensitivity may corrupt the first-pass retrieval and misdirect the expansion—has not been investigated. Additionally, the interaction between corpus-steered expansion and hybrid BM25–Dense fusion remains underexplored. While Exp4Fuse [?] shows that fusing the original- and expanded-query result lists from a *single sparse* retriever outperforms using the expansion alone, the asymmetric assignment of expansion across retriever *types* in a heterogeneous dense–sparse hybrid—and its behaviour for Arabic—has not been studied. Whether applying query expansion to only one retriever in such a hybrid can outperform applying it to both is therefore an open question with practical implications for retrieval pipeline design.

This thesis addresses these questions through a systematic evaluation of ten open-source language models for Query2Doc-based query expansion on the MIRACL Arabic benchmark, as described in the following chapter.

2.6 Chapter Summary

This chapter established the theoretical and empirical foundations for the research presented in this thesis. The key points are summarized as follows:

- RAG systems address the hallucination and knowledge staleness limitations of LLMs by incorporating external knowledge retrieval (Section 2.1.2), but their effectiveness is bounded by retrieval quality.
- Information retrieval can be performed via sparse methods (BM25), dense methods (DPR), or their combination, each with distinct strengths (Section 2.1.3).
- Query enhancement techniques—particularly LLM-based generative expansion—offer a modular approach to improving retrieval effectiveness without modifying the index or generator (Section 2.1.4).

- Arabic presents unique challenges for information retrieval due to morphological richness, diglossia, and orthographic variations (Section 2.1.5).
- The mathematical models for BM25 scoring, cosine similarity, the hybrid fusion functions, and the evaluation metrics Recall@ k , NDCG@ k , and MRR were formulated (Section 2.2).
- The MIRACL Arabic benchmark was selected for evaluation after surveying candidate Arabic retrieval and QA datasets against retrieval-suitability criteria (Section 2.3).
- Ten language models spanning Arabic-specialized and multilingual architectures were described for use in the comparative evaluation (Section 2.4).
- Recent literature demonstrates that 7–8B parameter models are sufficient for English query expansion, but no study has evaluated this approach for Arabic (Section 2.5).
- A clear research gap exists at the intersection of query enhancement techniques and Arabic information retrieval, which this thesis addresses through systematic model comparison on the MIRACL benchmark (Section 2.5.4).

The following chapter presents the methodology employed to address this research gap, describing the experimental design, dataset, evaluation pipeline, and the specific configurations used for each model.

Chapter 3

Methodology

This chapter presents the methodology employed throughout this research. The experimental work was conducted in a systematic, phased approach: establishing retrieval baselines (Section 3.2), performing error analysis to identify failure patterns (Section 3.3), implementing and adapting the Query2Doc query enhancement technique (Section 3.4), and conducting a comprehensive model comparison across ten open-source Large Language Models (Section 3.5). The chapter then presents the query repetition methodology for resolving BM25 term-dilution (Section 3.6), the hybrid BM25–Dense fusion methodology (Section 3.7), the Corpus-Steered Query Expansion (CSQE) pipeline including its component ablation design and retriever-specific application strategy (Section 3.8), and the per-query error analysis methodology (Section 3.9). All experimental setup details, including dataset characteristics and hardware configurations, are described in Section 3.1. The results corresponding to each methodological phase are presented in Chapter 4.

3.1 Dataset and Experimental Setup

3.1.1 Dataset: MIRACL Arabic

All experiments in this thesis were conducted on the Arabic subset of the Multilingual Information Retrieval Across a Continuum of Languages (MIRACL) benchmark [36]. MIRACL was selected as the primary dataset for three reasons: (1) it is retrieval-focused with human-annotated relevance judgements, (2) it contains natural query-document mismatch characteristic of real-world Arabic information needs, and (3) it is the standard benchmark for multilingual retrieval research, enabling comparison with published results.

The MIRACL Arabic corpus consists of 2,061,414 passages derived from 656,982 Arabic Wikipedia articles. Each passage is identified by a document identifier in the format **X#Y**, where **X** denotes the article and **Y** the passage number within that article. The development set, which was used for all experiments reported in this thesis, contains 2,896 queries with 29,197 human-annotated relevance judgements. Each judgement assigns a relevance grade ranging from 0 (not relevant) to 3 (highly relevant), provided by native Arabic speakers. The corpus is written exclusively in Modern Standard Arabic (MSA); consequently, the evaluation of dialectal Arabic queries falls outside the scope of this work.

The dataset was accessed directly from HuggingFace¹ without local download, as both the corpus and pre-built retrieval indices were available through the Pyserini toolkit [37].

3.1.2 Hardware and Software Environment

All experiments were conducted on Google Colab, utilising two GPU configurations depending on availability and model requirements:

- **NVIDIA T4 (15 GB VRAM):** Used for baseline experiments and models with fewer than 4 billion parameters in FP16 precision.
- **NVIDIA A100 (40 GB VRAM):** Used for models requiring 4-bit quantisation or exceeding 7 billion parameters, and for batch sizes larger than 8.

The software environment consisted of Python 3.10, PyTorch 2.0+, the Hugging Face Transformers library (v4.30+), Pyserini (v0.22.1) for index access, and FAISS for dense vector search. For models exceeding GPU memory in native precision, 4-bit NormalFloat (NF4) quantisation was applied using the `bitsandbytes` library [38], following the QLoRA quantisation scheme.

¹<https://huggingface.co/datasets/miracl/miracl>

3.1.3 Evaluation Metrics

Three standard information retrieval metrics were employed for evaluation, as defined in Section 2.2:

- **NDCG@10** (Equation 2.9): Measures the quality of ranking within the top-10 retrieved passages. This was treated as the primary metric throughout all experiments.
- **Recall@10** (Equation 2.7): Measures the proportion of all relevant passages that appear in the top-10 results.
- **MRR** (Equation 2.10): Measures how quickly the first relevant passage is retrieved.

Additionally, Recall@100 was computed to assess overall retrieval coverage. All metrics were calculated using the `pytrec_eval` library over the full development set of 2,896 queries.

3.1.4 Experimental Pipeline Overview

The experimental pipeline follows a two-phase architecture, as illustrated in Figure 3.1. In the first phase (query enhancement), each query from the MIRACL development set is processed by a query enhancement module, which may leave the query unchanged (baseline) or augment it using an LLM-generated pseudo-document. In the second phase (retrieval and evaluation), the enhanced query is passed to a retrieval system, and the top-100 results are evaluated against the MIRACL relevance judgements.




Figure 3.1: Overview of the experimental pipeline. A user query passes through the query enhancement layer; BM25S and mDPR retrieve independently; their ranked lists are fused (CC or RRF) and the top- k passages are scored against the MIRACL Arabic dev qrels. The query enhancement module is the variable component across experiments; the retrieval and evaluation components remain fixed.

A key design decision was to separate query generation from retrieval evaluation into two independent notebooks. The query enhancement phase (which requires GPU for LLM inference) saves enhanced queries to disk as serialised Python objects. The evaluation phase then loads these pre-generated queries and performs retrieval and metric computation. This two-notebook workflow enabled efficient reuse of enhanced queries across multiple retrieval configurations (Dense and BM25) without regenerating them.

3.2 Baseline Implementation

Two retrieval baselines were established to provide reference performance levels against which all query enhancement experiments would be compared: a dense retrieval baseline using mDPR and a sparse retrieval baseline using BM25S. Both baselines used an `IdentityEnhancer` that returned queries unchanged.

3.2.1 Dense Retrieval Baseline (mDPR)

The dense retrieval baseline employed Multilingual Dense Passage Retrieval (mDPR), described in Section 2.4.3.1. The mDPR encoder (`castorini/mdpr-tied-pft-msmarco`) was loaded on GPU, and queries were encoded in batches of 64. The pre-built FAISS index

for MIRACL Arabic (5.47 GB, 768-dimensional embeddings) was accessed via Pyserini. For each query, the top-100 passages were retrieved using cosine similarity.

mDPR was pre-trained on MS MARCO (English) and was *not* fine-tuned on the MIRACL Arabic corpus, unlike models such as BGE-M3 and E5 which were partially trained on MIRACL. This choice mirrors a realistic deployment scenario in which an off-the-shelf multilingual retriever is improved through query-side interventions rather than corpus-specific fine-tuning, and ensures that observed improvements reflect query enhancement rather than retriever memorisation of evaluation queries. Figure 3.2 illustrates the mDPR query-encoding pipeline.




Figure 3.2: Dense retrieval pipeline using multilingual DPR. Queries are encoded in batches on a single A100 GPU under FP16; the resulting 768-dimensional vectors are searched against a pre-built FAISS index of the corpus.

3.2.2 Sparse Retrieval Baseline (BM25S)

The sparse retrieval baseline employed BM25S [23], a pure-Python implementation of the BM25 algorithm (described in Section 2.1.3.1). BM25S was selected over the Pyserini-based BM25 implementation after encountering a blocking dependency conflict between Java 21 and Java 11 in the Google Colab environment. The BM25S library eliminates all Java dependencies while maintaining competitive performance.

The BM25S index was built from the MIRACL Arabic corpus with the following configuration:

- **Tokenisation:** Arabic stemming using PyStemmer with NLTK Arabic stopwords (245+ words).
- **BM25 variant:** Lucene-style scoring with parameters $k_1 = 0.9$ and $b = 0.4$.
- **Index storage:** Google Drive (~ 0.5 GB), loaded via symbolic links in Colab.

This configuration achieved 96% of the official MIRACL Pyserini baseline performance, with the 4% difference deemed acceptable given the benefits of a pure-Python pipeline that integrates directly with the query enhancement modules. Figure 3.3 summarises the BM25S indexing pipeline.




Figure 3.3: BM25S indexing pipeline applied to the MIRACL Arabic corpus. Passages are tokenised with PyStemmer, stopwords-removed using 245 Arabic stopwords, then indexed under BM25S with $k_1 = 0.9$ and $b = 0.4$.

3.2.3 Baseline Comparison Rationale

Dense and sparse retrieval were tested *separately* rather than combined into a hybrid system. This decision, confirmed during the project planning phase, was made to provide

clearer insight into where query enhancement improvements originate. As discussed in Section 2.1.3, dense retrieval captures semantic similarity while sparse retrieval relies on exact term matching. Testing both baselines independently enables analysis of whether a given query enhancement technique benefits semantic matching, lexical matching, or both.

3.3 Error Analysis Methodology

Before implementing query enhancement techniques, a systematic error analysis was conducted on the dense retrieval baseline to identify the primary failure patterns and guide technique selection.

3.3.1 Quantitative Analysis Framework

A quantitative analysis was performed over all 2,896 queries in the MIRACL Arabic development set. Queries were segmented into three performance categories based on their individual NDCG@10 scores:

- **Failed:** $\text{NDCG@10} < 0.3$
- **Mediocre:** $0.3 \leq \text{NDCG@10} < 0.7$
- **Successful:** $\text{NDCG@10} \geq 0.7$

These three categories describe each query’s *absolute* retrieval quality in isolation. They should be distinguished from the *pairwise* classification scheme introduced later for the corpus-steered query expansion analysis (Section 3.9), which measures the *change* in NDCG@10 between two systems rather than absolute success, together with the regression sub-typing built upon it. Each scheme answers a different question—an absolute success rating versus the magnitude and direction of a system-to-system change. The 0.3 and 0.7 cut-offs adopted here are pragmatic boundaries chosen for interpretability rather than derived from a formal statistical criterion.

For each query, the following features were computed:

- a. **Query length** (in words), measured after whitespace tokenisation.
- b. **Score gap** between the top-ranked passage and the second-ranked passage, as a proxy for retriever confidence.
- c. **First relevant document rank**, indicating how deep the retriever must search before finding a relevant passage.
- d. **Coverage at multiple cut-offs** (top-10, top-20, top-50, top-100), measuring the proportion of queries with at least one relevant passage at each depth.

The Pearson correlation coefficient was computed between query length and NDCG@10 to quantify the relationship between query verbosity and retrieval performance.

3.3.2 Query Length Bucketing

Queries were grouped into three length buckets to analyse the effect of query length on retrieval performance:

- **Short:** 1–3 words (5.1% of queries)
- **Medium:** 4–8 words (86.2% of queries)
- **Long:** 9+ words (8.8% of queries)

The mean NDCG@10 was computed for each bucket, and the performance gap between short and long queries was used to quantify the impact of information poverty on retrieval effectiveness.

3.3.3 Failed Query Inspection

The 20 worst-performing queries (NDCG@10 = 0.000, indicating zero relevant passages in the top-10) were manually inspected to identify recurring linguistic patterns. Each query

was categorised by question type (factoid, definition, temporal), the presence of named entities, diacritics, technical terminology, and spelling variations. These observations informed the selection of the first query enhancement technique.

3.4 Query2Doc Implementation

Based on the error analysis findings (presented in Section 4.2), Query2Doc [6] was selected as the first query enhancement technique. This section describes the Query2Doc approach, the modifications made for Arabic zero-shot application, and the implementation details.

3.4.1 Query2Doc Technique

Query2Doc, introduced by Wang et al. [6], uses an LLM to generate a pseudo-document—a short passage that attempts to answer the query. This pseudo-document is then concatenated with the original query to form an enhanced query. The rationale is that the pseudo-document introduces additional vocabulary, context, and semantic information that bridges the gap between the query and relevant passages.

As described in Section 2.1.4, Query2Doc falls within the category of LLM-based generative query enhancement techniques. It differs from Hypothetical Document Embeddings (HyDE) [8] in a critical way: while HyDE uses the generated document *alone* as the retrieval query, Query2Doc *concatenates* the generated document with the original query, preserving the original query terms. The Query2Doc generation pipeline used in this thesis is illustrated in Figure 3.4.




Figure 3.4: Query2Doc generation pipeline: the original query is wrapped in a fixed prompt template; an open-source LLM generates a short pseudo-document; the pseudo-document is concatenated with the query before being passed to the retriever.

3.4.2 Modifications from the Original Paper

Several modifications were made to adapt Query2Doc for this research context:

- a. **Zero-shot prompting:** The original Query2Doc paper employed few-shot prompting with 4 demonstration examples using GPT-3 (text-davinci-003, 175B parameters). In this work, zero-shot prompting was adopted exclusively. This decision was driven by resource constraints (small open-source models have limited context windows for few-shot examples) and the desire to evaluate model capabilities without task-specific demonstrations.
- b. **Small open-source models:** Rather than using proprietary API-based models, all experiments utilised open-source models with 2–8 billion parameters that can run on Google Colab free-tier GPUs. This approach ensures reproducibility and eliminates API costs.

- c. **Arabic-only generation:** The system prompt explicitly instructs the model to respond in Arabic only, whereas the original paper targeted English queries.
- d. **Retriever-specific concatenation:** For dense retrieval, simple concatenation was used: $q_{\text{enhanced}} = q \oplus d_{\text{pseudo}}$, where $\oplus$ denotes string concatenation. The original paper recommends repeating the query $n = 5$ times before concatenation for sparse retrieval to prevent term dilution, as discussed in Section 3.4.6.

3.4.3 LLM Configuration

The initial Query2Doc experiment used Qwen 2.5 3B Instruct (described in Section 2.4.2.1) as the generation model. The following system prompt was used across all experiments:

“You are asked to write a passage that answers the given query. Do not ask the user for further clarification. Respond in Arabic only.”

The generation parameters were configured as follows:

Table 3.1: LLM generation parameters for the initial Query2Doc experiment.

Parameter	Value
Maximum new tokens	128
Temperature	0.7
Top- p (nucleus sampling)	0.9
Batch size	8
Precision	FP16

The maximum generation length was set to 128 tokens (~ 100 Arabic words), providing sufficient context for generating informative pseudo-documents while keeping per-query generation time practical.

3.4.4 Batch Processing Optimisation

A critical engineering challenge was generation throughput. Naïve sequential processing (~ 10 seconds per query) would require approximately 8 hours for the full development set. Three optimisations were implemented to achieve practical runtimes:

- a. **Batch processing:** Multiple queries were tokenised with left-padding and processed simultaneously through the LLM, substantially reducing per-query generation time.
- b. **Reduced token generation:** The maximum generation length was set to 128 tokens, keeping generation time practical while retaining expansion quality.
- c. **Inference optimisations:** The model was set to evaluation mode (`model.eval()`), half-precision (`float16`) was used, and gradient computation was disabled (`torch.no_grad()`).

Together, these optimisations reduced full-set generation time to approximately 40 minutes per model on a T4 GPU, enabling practical iteration across multiple models and configurations.

3.4.5 Dense Retrieval with Query2Doc

For dense retrieval experiments, the enhanced query was formed by simple concatenation:

$$q_{\text{enhanced}} = q_{\text{original}} \oplus d_{\text{pseudo}} \quad (3.1)$$

The enhanced query was then encoded by the mDPR encoder and used for FAISS similarity search, following the same procedure as the baseline (Section 3.2.1). This approach leverages the semantic content of the pseudo-document to improve the query's embedding representation.

3.4.6 BM25 Retrieval with Query2Doc

For sparse retrieval, the same enhanced queries were evaluated against the BM25S index. Notably, the initial BM25 experiment used simple concatenation identical to the dense configuration, *without* implementing the query repetition strategy recommended by Wang et al. [6]. The original paper recommends:

$$q_{\text{enhanced}}^{\text{BM25}} = \underbrace{q \oplus q \oplus \dots \oplus q}_{n \text{ times}} \oplus d_{\text{pseudo}} \quad (3.2)$$

where $n = 5$. This repetition preserves the importance of original query terms in the BM25 scoring function (Equation 2.2), preventing the pseudo-document from diluting the term weights of the original query. The omission of this strategy in the initial experiment, and its subsequent impact, is analysed in Section 4.3.2.

3.5 Model Comparison Methodology

Following the initial Query2Doc results with Qwen 2.5 3B, a comprehensive model comparison was conducted to identify the optimal LLM for Arabic query expansion. Ten open-source models were evaluated using the identical Query2Doc pipeline.

3.5.1 Model Selection Criteria

Models were selected based on three criteria: (1) Arabic language support, either through dedicated Arabic training or strong multilingual capabilities; (2) parameter count within the 2–8 billion range to remain executable on Google Colab GPUs; and (3) open-source availability with permissive licences for research use. The selected models span two categories, both of which are described in detail in Section 2.4:

- **Arabic-specialised:** Falcon-H1-3B (Section 2.4.1.1), Jais-2-8B (Section 2.4.1.2), ALLaM-7B (Section 2.4.1.3), and SILMA Kashif-2B (Section 2.4.1.4).
- **Multilingual:** Qwen 2.5 3B (Section 2.4.2.1), Qwen 2.5 7B (Section 2.4.2.2), Qwen3-4B (Section 2.4.2.3), Qwen3-8B (Section 2.4.2.4), Gemma 3 4B (Section 2.4.2.5), and Aya Expanse 8B (Section 2.4.2.6).

Table 3.2 summarises the configuration used for each model.

Table 3.2: Model configurations for the Query2Doc comparison experiments. All models used `max_new_tokens = 128` and `top_p = 0.9` unless otherwise noted.

Model	Params	Precision	Temp.	Batch	GPU
Qwen 2.5 3B	3B	FP16	0.7	8	T4
SILMA 2B	2B	FP16	0.1	16	A100
Falcon-H1 3B	3.15B	BF16	0.1	1*	A100
Qwen3-4B	4B	FP16	0.7	32	A100
Gemma 3 4B	4B	FP16	0.1	16	A100
Qwen 2.5 7B	7B	FP16	0.1	16	A100
ALLaM 7B	7B	FP16	0.7	—	A100
Jais-2 8B	8B	BF16	0.1	8	A100
Qwen3-8B	8B	FP16	0.1	16	A100
Aya Expansive 8B	8B	NF4	0.1	8	A100

*Forced to batch size 1 due to architecture-specific batching bug (see Section 3.5.4).

3.5.2 Standardised Evaluation Protocol

To ensure fair comparison, all models were evaluated using an identical protocol:

- a. **Sanity check:** The first 5 queries were processed and the outputs were manually inspected to verify that the model (a) generates text in Arabic, (b) produces relevant content, and (c) does not exhibit degenerate behaviour (repetition, code output, or empty responses).
- b. **Full generation:** All 2,896 queries were processed using the Query2Doc pipeline, and the enhanced queries were saved to disk.
- c. **Dense evaluation:** Enhanced queries were evaluated against the mDPR dense retrieval baseline.
- d. **BM25 evaluation:** The same enhanced queries were evaluated against the BM25S sparse retrieval baseline.

The work was divided between the two team members: one researcher tested Falcon-H1-3B, Jais-2-8B, ALLaM-7B, and Qwen3-4B, while the other tested SILMA Kashif-2B, Qwen 2.5-7B, Qwen3-8B, Gemma 3 4B, and Aya Expansive 8B. All experiments used the same codebase, system prompt, and evaluation pipeline to minimise confounding variables.

3.5.3 Temperature Selection

Temperature was not fixed at a single value across all models. Two considerations guided temperature selection:

- a. **Model-recommended values:** Some models specify recommended generation parameters. Falcon-H1 recommends temperature 0.1, noting that higher values “largely drop performance.” Qwen3 models require sampling (`do_sample=True`) to avoid infinite repetitions with greedy decoding.
- b. **Empirical testing:** For SILMA Kashif-2B, two temperature values were compared (0.7 and 0.1). Temperature 0.1 yielded a 2.5% improvement in NDCG@10 over 0.7, leading to its adoption as the default for subsequent experiments where no model-specific recommendation existed.

3.5.4 Model-Specific Technical Issues

Several models required specific engineering workarounds, documented here for reproducibility:

- a. **Falcon-H1-3B:** The hybrid Mamba2-Transformer architecture (described in Section 2.4.1.1) contains a batching bug in the attention mask extension logic. When left-padded batches are used, the 4D causal attention mask is not extended as new tokens are generated, causing a shape mismatch. The workaround was to force `batch_size=1`, resulting in significantly longer generation times but correct outputs.
- b. **Jais-2-8B:** This model requires BF16 precision rather than FP16 due to its Squared-ReLU activation function, which produces values that overflow FP16 range. Additionally, the `token_type_ids` parameter must be excluded from the model inputs, and the padding token must be explicitly set to the end-of-sequence token.
- c. **Qwen3-4B and Qwen3-8B:** The Qwen3 generation includes “thinking” tokens (internal chain-of-thought reasoning) that must be stripped from the output. The parameter `enable_thinking=False` was passed during generation. Greedy decoding must not be used, as it causes infinite repetition loops.

- d. **ALLaM-7B:** As a preview-stage model, ALLaM exhibited a sentencepiece tokenizer artefact where the Unicode character “ ” (lower one-eighth block) leaked into decoded outputs. This special character, used internally by the sentencepiece tokenizer to denote word boundaries, contaminated the pseudo-documents. The impact on retrieval performance is discussed in Section 4.4.3.

3.5.5 Quantisation Strategy

Models exceeding the available GPU memory in native precision were quantised using 4-bit NormalFloat (NF4) quantisation via the `bitsandbytes` library [38]. The quantisation configuration was:

- **Quantisation type:** NF4 (NormalFloat 4-bit)
- **Compute dtype:** BF16 (for matrix multiplication)
- **Double quantisation:** Enabled (quantises the quantisation constants for further memory savings)

This configuration was applied to Jais-2-8B (on T4 only; BF16 was used on A100) and Aya Expanse 8B (ALLaM-7B was run in full FP16 precision). Notably, Aya Expanse 8B achieved the best overall performance despite being quantised to 4-bit, suggesting that model architecture and training data quality are more significant factors than numerical precision for this task.

3.6 Query Repetition for Sparse Retrieval

The model comparison experiments (Section 4.4.2) revealed that single-pass Query2Doc ($n = 1$) degraded BM25 performance for six of nine models due to term dilution (Section 4.3.2.1). The original Query2Doc paper [6] recommended concatenating $n = 5$ pseudo-documents with equal weighting, but did not investigate the sensitivity of the repetition factor to

model size or pseudo-document length. Two solution families were evaluated to address this limitation.

Fixed Repetition (Query2Doc-style). The original query is prepended n times before the pseudo-document, where n is a global hyperparameter applied identically to every query regardless of its length or the pseudo-document length. The enhanced query is assembled as:

$$q_{\text{rep}} = \underbrace{q \circ q \circ \dots \circ q}_{n \text{ times}} \circ d_1 \circ d_2 \circ \dots \circ d_k \quad (3.3)$$

where $\circ$ denotes string concatenation with a space separator and k is the number of pseudo-documents (for single-pass Query2Doc, $k = 1$). Fixed repetition was swept over $n \in \{1, 3, 5, 7, 10\}$.

Adaptive Repetition (MuGI-style). Zhang et al. [26] proposed computing n per-query from the ratio of pseudo-document length to query length, controlled by a parameter β :

$$n(q, d, \beta) = \max\left(1, \left\lfloor \frac{|d|}{|q| \cdot \beta} \right\rfloor\right) \quad (3.4)$$

where $|q|$ and $|d|$ are the token lengths of the query and pseudo-document respectively. Larger β yields fewer repetitions; smaller β yields more. Adaptive repetition was swept over $\beta \in \{2, 4, 6\}$.

Motivation for the Adaptive Variant. The optimal number of repetitions depends on the length ratio between query and pseudo-document. A 3-word query paired with a 200-token pseudo-document requires more repetition than a 15-word query paired with a 100-token pseudo-document. The MuGI per-query adaptation removes this length sensitivity and, as reported in Section 4.6, produced the best results for the two largest models (Aya Expansive 8B and Jais-2 8B), where $\beta = 2$ was the winning configuration.

Sweep Design. Nine models were evaluated across eight configurations ($n \in \{1, 3, 5, 7, 10\}$ and $\beta \in \{2, 4, 6\}$), yielding 72 BM25 evaluations. All pseudo-document expansions had been generated during the model comparison phase (Section 3.5); no new LLM inference was required. The stored expansion files (serialised Python objects) were loaded and combined with varying repetition factors, then evaluated using `pytrec_eval` on the full MIRACL Arabic development set (2,896 queries).

3.7 Hybrid Retrieval Fusion

The baseline experiments (Section 4.1.1) confirmed that BM25 and mDPR exhibit complementary strengths: mDPR achieves superior ranking quality while BM25 provides broader lexical coverage. Hybrid retrieval fusion combines the ranked lists from both retrievers into a single ranking to exploit this complementarity. Two fusion methods, introduced in Section 2.1.3.3, were evaluated.

Reciprocal Rank Fusion (RRF). RRF (Equation 2.5) combines ranked lists by summing inverse rank scores with a smoothing constant k . Two values were tested: $k = 20$, following the standard practice established by Bruch et al. [7], and $k = 60$.

Convex Combination (CC). CC (Equation 2.6) computes a weighted sum of normalised retriever scores, where α controls the relative contribution of the dense retriever. Before combination, the scores from each retriever were min-max normalised to $[0, 1]$ on a per-query basis to account for the different score scales of BM25 and mDPR. The weight α was swept over $\{0.1, 0.2, \dots, 0.9\}$.

Setup. Both BM25S and mDPR retrieved the top-100 candidates independently using the original (unenhanced) queries. Fusion and re-ranking were performed in Python without additional GPU inference. The resulting hybrid baseline establishes the performance ceiling that all subsequent query enhancement methods must surpass. The two fusion methods are illustrated side by side in Figure 3.5.




Figure 3.5: Two retriever-fusion methods evaluated in this thesis. Convex Combination (left) computes a weighted sum of normalised retrieval scores; Reciprocal Rank Fusion (right) sums the reciprocal of each retriever’s rank for every document.

3.8 Corpus-Steered Query Expansion (CSQE)

Corpus-Steered Query Expansion (CSQE), introduced by Lei et al. [25] and described in Section 2.5.2, extends blind Query2Doc by grounding LLM-generated expansions in first-pass retrieved documents. This section describes the CSQE pipeline as adapted for Arabic retrieval, the component ablation design, and the retriever-specific application strategy.

3.8.1 The CSQE Pipeline

The CSQE pipeline distinguishes itself from standard (blind) Query2Doc through a two-stage design:

- a. **First-pass retrieval.** BM25S retrieves the top- $k = 5$ documents using the original query q .
- b. **Corpus-grounded expansion.** Each retrieved document d_i is concatenated with q and passed to the LLM with a CSQE-specific system prompt instructing the extraction of topically relevant Arabic vocabulary grounded in the retrieved passage.

- c. **Blind expansion.** The same query q is passed to the LLM without any retrieved context, following the standard Query2Doc procedure.
- d. **Expansion assembly.** The corpus-grounded and blind samples are combined with a repeated original query:

$$q_{\text{CSQE}} = \underbrace{q \circ q \circ \dots \circ q}_{\alpha \text{ times}} \circ c_1 \circ c_2 \circ b_1 \circ b_2 \quad (3.5)$$

where c_i denotes corpus-grounded expansion samples, b_i denotes blind expansion samples, α is the query repetition factor (Equation 3.3), and each expansion sample appears exactly once ($k = 1$ in MuGI notation).

The configuration used for the primary CSQE experiment was: $k = 5$ first-pass documents, 2 corpus-grounded samples, 2 blind samples, $\alpha = 4$, temperature 1.0, `max_new_tokens` = 128 per sample, using Aya Expanse 8B (Section 2.4.2.6) in BF16 precision on an A100 GPU. Figure 3.6 illustrates the three-stage CSQE pipeline alongside a worked example.



Figure 3.6: Corpus-Steered Query Expansion (CSQE) pipeline. Stage 1: BM25 retrieves the top- $k_1 = 5$ first-pass passages. Stage 2: Aya Expanse 8B is invoked four times ($\alpha = 4$, temp = 1.0): two corpus-grounded calls (c_1, c_2) extract key sentences from the retrieved passages; two blind calls (b_1, b_2) generate from parametric knowledge. Stage 3: the original query (repeated α times) is concatenated with all four pseudo-documents and fed to a BM25 second pass, yielding $\text{nDCG}@10 = 0.616$.

CSQE System Prompt. For corpus-grounded generation, the model was given a system prompt distinct from the blind Query2Doc prompt of Section 3.4.3, instructing it to *extract* content grounded in the retrieved passages rather than to generate freely:

“You are an information retrieval assistant. You will examine retrieved documents and extract key sentences relevant to the query. The documents are in Arabic. Respond in Arabic only.”

The accompanying user message presented the original query followed by the numbered first-pass documents and a closing instruction to “begin by examining the initially retrieved documents and identifying the ones that are relevant, even partially, to the query ... then extract the key sentences from each document that contribute to their relevance.” A single English worked example (a query about warm-blooded sharks, taken from the original CSQE paper [25]) was prepended as a one-shot demonstration; Aya Expanse, being multilingual, transfers this pattern to Arabic queries while still producing Arabic output. The blind-expansion branch reused the Query2Doc system prompt of Section 3.4.3 verbatim.

Rationale for Combining Corpus and Blind Samples. Corpus-grounded samples provide attested Arabic vocabulary anchored to the actual Wikipedia passages retrieved in the first pass. Blind samples provide diverse answer-space vocabulary unconstrained by the corpus. The two sources are hypothesised to be complementary: corpus samples prevent hallucination of incorrect entities, while blind samples expand vocabulary coverage beyond the retrieved passages. This hypothesis is validated through the ablation study presented in Section 4.8.2.

3.8.2 Component Ablation Design

To isolate the individual contributions of corpus-grounded and blind expansion, three ablation variants were defined:

- **Corpus-only ablation:** `num_corpus_samples = 4, num_blind_samples = 0.`

This variant isolates the contribution of first-pass grounding by using only corpus-grounded expansions.

- **Blind-only ablation:** `num_corpus_samples = 0`, `num_blind_samples = 4`. This variant isolates the contribution of blind generation, equivalent to a 4-sample Query2Doc with repetition.
- **Full CSQE configuration:** `num_corpus_samples = 2`, `num_blind_samples = 2`. The combined system as described above.

All three variants used $\alpha = 4$ and the same Aya Expanse 8B model. Additionally, a query repetition factor sweep ($\alpha \in \{1, 2, 3, 4\}$) was conducted on the full CSQE variant. This sweep was reconstructed from stored expansion files (serialised Python objects) without requiring new LLM inference—only the number of query prepends was varied before re-running BM25 evaluation.

3.8.3 Retriever-Specific Application

The CSQE-expanded query (q_{CSQE}) is substantially longer than the original query. Three fusion configurations were tested to determine the optimal retriever–query assignment; whether expansion helps or hurts each retriever was left as an empirical question, with the causal interpretation reserved for the results discussion (Section 4.9.1).

- **BM25-expanded:** BM25 receives the CSQE-expanded query; Dense receives the original query.
- **Dense-expanded:** BM25 receives the original query; Dense receives the CSQE-expanded query.
- **Both-expanded:** Both BM25 and Dense receive the CSQE-expanded query.

For each configuration, both RRF ($k = 20$) and CC (α swept over $\{0.1, \dots, 0.9\}$) were evaluated. This design isolates the effect of retriever-specific query representation

and determines whether asymmetric expansion outperforms symmetric application. The best-performing configuration, identified empirically in Chapter 4, is illustrated in Figure 3.7.




Figure 3.7: Architecture of the thesis’s best-performing system. CSQE-expanded queries are fed to BM25 while the original short queries are fed to mDPR; the two ranked lists are fused with Reciprocal Rank Fusion ($k = 20$), yielding the headline nDCG@10 of 0.714.

3.9 Per-Query Error Analysis

Aggregate metrics such as NDCG@10 can mask substantial per-query variation: a system with a high average may still regress on a significant minority of queries. A per-query error analysis was conducted to decompose the overall CSQE improvement and identify the conditions under which corpus-steered expansion helps or hinders retrieval.

Per-Query Metric Computation. The `pytrec_eval` library was used to compute NDCG@10 for each of the 2,896 MIRACL Arabic development queries under both the CSQE+Hybrid system and the Aya Blind BM25 $n = 1$ baseline. The per-query delta was defined as $\Delta_i = \text{NDCG@10}_{\text{CSQE},i} - \text{NDCG@10}_{\text{Blind},i}$.

Classification Thresholds. Three categories were defined based on the per-query scores and deltas:

- **Failure:** CSQE NDCG@10 < 0.1 . Queries where the system fails to retrieve any relevant document in the top-10, regardless of the baseline performance.
- **Big win:** $\Delta_i > 0.3$. Queries where CSQE achieves a large improvement over the blind baseline.

- **Regression:** $\Delta_i < -0.1$. Queries where CSQE performs worse than the blind baseline by a non-trivial margin.

First-Pass Quality Split. To assess the role of BM25 first-pass quality in CSQE effectiveness, queries were split into two groups based on whether the BM25 first-pass retrieval returned a relevant document at *rank 1* (a positive relevance judgement, $\text{rel} > 0$, for the top-1 retrieved document in the MIRACL qrels). Although CSQE grounding draws on the top- $k = 5$ first-pass documents, the rank-1 result is used as the indicator of first-pass quality because it most directly reflects whether the retriever has correctly identified the query topic. This split enables analysis of whether CSQE gains are driven primarily by successful corpus grounding or by the blind expansion component.

Regression Classification. Regression queries ($\Delta_i < -0.1$) were manually inspected and classified into three types:

- **Type A (Strong BM25 hurt):** The BM25 baseline already achieved $\text{NDCG@10} \geq 0.3$ for the query; the CSQE expansion introduced off-topic vocabulary that diluted the effective query terms.
- **Type B (Poisoned first-pass):** The BM25 baseline achieved $\text{NDCG@10} < 0.1$ for the query; the first-pass retrieval returned irrelevant documents, causing the LLM to ground its expansion on incorrect content.
- **Type C (Partial BM25):** The BM25 baseline achieved NDCG@10 between 0.1 and 0.3; a mixed-quality scenario where the expansion may have helped or hurt depending on specific query characteristics.

Chapter 4

Results and Discussion

This chapter presents the results of all experiments conducted in this research, following the zigzag structure with the methodology described in Chapter 3. Section 4.1 reports the baseline retrieval performance. Section 4.2 presents the error analysis findings that guided technique selection. Section 4.3 reports the initial Query2Doc results with both dense and sparse retrieval. Section 4.4 presents the comprehensive model comparison across ten LLMs. Section 4.5 synthesises the key findings from the model comparison phase. The chapter then reports the expanded experimental results: the BM25 query repetition sweep (Section 4.6), hybrid retrieval fusion (Section 4.7), Corpus-Steered Query Expansion (Section 4.8), CSQE combined with hybrid fusion (Section 4.9), and a per-query error analysis decomposing the sources of improvement and regression (Section 4.10).

4.1 Baseline Results and Comparison

The two retrieval baselines described in Section 3.2 were evaluated on the full MIRACL Arabic development set (2,896 queries). Table 4.1 presents the results.

Table 4.1: Baseline retrieval results on MIRACL Arabic dev set (2,896 queries). Bold marks the higher score between the two single-retriever baselines; the no-QE hybrid fusion result is included as the strongest non-enhancement reference (detailed in Section 4.7).

Retriever	NDCG@10	Recall@10	Recall@100	MRR
mDPR (Dense)	0.4993	0.6156	0.8407	0.5328
BM25S (Sparse)	0.4621	0.5964	0.8577	0.4836
Hybrid RRF ($k = 20$, no QE)	0.6267	0.7597	0.9466	0.6517

The mDPR dense baseline reproduced the official MIRACL result exactly: this work measured $\text{NDCG@10} = 0.4993$, and the published value of 0.499 is the same number

rounded to three decimal places. This confirmed the correctness of the implementation. The BM25S sparse baseline achieved 96% of the official Pyserini BM25 performance ($\text{NDCG@10} = 0.4621$ vs. target 0.481), with the 4% difference attributable to the use of BM25S rather than Pyserini’s Java-based implementation, as discussed in Section 3.2.2.

4.1.1 Complementary Strengths

The two baselines exhibited complementary performance characteristics. mDPR achieved higher NDCG@10 (+8.1%) and MRR (+10.2%), indicating superior ranking quality—relevant passages were placed higher in the result list. BM25S achieved higher Recall@100 (+2.0%), indicating broader retrieval coverage—more relevant passages were retrieved overall, though ranked less effectively.

This complementary behaviour aligns with the theoretical properties of dense and sparse retrieval discussed in Section 2.1.3: dense retrieval captures semantic similarity and ranks well, while sparse retrieval excels at broad lexical matching. This finding confirmed the decision to test query enhancement techniques on both retrievers independently. It also motivates a third reference point: fusing the two retrievers without any query enhancement (Reciprocal Rank Fusion, $k = 20$) already raises NDCG@10 to 0.6267 (Table 4.1). This no-QE hybrid is reported here, alongside the single-retriever baselines, as the strongest non-enhancement reference; it is analysed in detail in Section 4.7, where it serves as the ceiling that all query enhancement methods must surpass. The per-query NDCG@10 distribution for both single-retriever baselines is shown in Figure 4.1.




Figure 4.1: Per-query NDCG@10 distribution for the two baseline retrievers on the MIRACL Arabic dev set. Both exhibit a bimodal distribution: a large mass near zero (poorly served queries) and another mass near one (perfectly served), with relatively few in between.

4.2 Error Analysis Findings

The error analysis methodology described in Section 3.3 was applied to the dense baseline. The findings guided the selection of the first query enhancement technique.

4.2.1 Overall Failure Rate

Table 4.2 presents the performance segmentation of all 2,896 queries.

Table 4.2: Performance segmentation of baseline queries by NDCG@10 score.

Category	Count	% of Total	Avg NDCG@10
Failed ($\text{NDCG@10} < 0.3$)	982	33.9%	0.053
Mediocre ($0.3 \leq \text{NDCG@10} < 0.7$)	962	33.2%	0.511
Successful ($\text{NDCG@10} \geq 0.7$)	952	32.9%	0.948

A notable finding was that 34% of queries (982 out of 2,896) were classified as failed, of which 736 retrieved no relevant passage at all within the top-10. The remaining queries split almost evenly between the mediocre (33%) and successful (33%) bands. While the substantial failed and mediocre fractions establish a clear opportunity for query enhancement, the analysis also shows that roughly a third of queries were already well served by the dense baseline. Figure 4.2 visualises this distribution on the cumulative per-query NDCG@10 curve.




Figure 4.2: Cumulative distribution of per-query NDCG@10 on the mDPR baseline. The annotated threshold at 0.3 separates the failed queries from the remainder.

4.2.2 Short Query Performance Gap

The analysis of retrieval performance by query length revealed a significant performance gap, as shown in Table 4.3.

Table 4.3: Retrieval performance by query length bucket.

Length Bucket	Avg NDCG@10	Count	% of Total
Short (1–3 words)	0.345	147	5.1%
Medium (4–8 words)	0.511	2,495	86.2%
Long (9+ words)	0.476	254	8.8%

Short queries achieved only 72% of the performance of long queries ($\text{NDCG@10} = 0.345$ vs. 0.476), a 28% gap. Performance did not, however, rise monotonically with length: medium-length queries (0.511) outscored long queries, and the Pearson correlation between query length and NDCG@10 was negligible ($r \approx -0.01$), indicating no linear length–performance trend. The salient pattern is therefore not that longer queries are better, but that the *shortest* queries specifically underperform—consistent with information poverty in queries of only one to three words. This motivated the selection of query expansion as the first enhancement technique: by generating pseudo-documents that add vocabulary and context, the effective information content of short queries is increased. The per-bucket distribution is shown in Figure 4.3.




Figure 4.3: Distribution of mDPR NDCG@10 within each query-length bucket. Short queries clearly underperform, motivating query enhancement as a remedy for information poverty in short queries.

4.2.3 Retrieval Coverage Analysis

The coverage analysis revealed an important distinction between retrieval and ranking:

Even at depth 100, the dense retriever failed to surface any relevant passage for roughly 10% of queries (coverage 90.1%), and for a further $\sim 15\%$ it retrieved a relevant passage but ranked it outside the top-10 (coverage 74.6% at depth 10). The first figure reflects

Table 4.4: Proportion of queries with at least one relevant passage at each retrieval depth.

Retrieval Depth	Coverage
Top-10	74.6%
Top-20	80.8%
Top-50	86.9%
Top-100	90.1%

a genuine *recall ceiling*—relevant content the dense encoder never retrieves—while the second reflects a ranking problem. Query enhancement can address both: enriching the query representation improves ranking, while combining the dense retriever with a lexical retriever (hybrid fusion, Section 4.7) directly raises the recall ceiling. The full Recall@ k curve for both baselines is plotted in Figure 4.4.




Figure 4.4: Mean Recall@ k for mDPR and BM25 on the MIRACL Arabic dev set, k on a log axis.

4.2.4 Technique Selection Rationale

Manual inspection of the 20 worst-performing queries revealed recurring failure patterns: vocabulary mismatch (e.g., the Arabic term “أزوت” for nitrogen versus the more common “نيتروجين”), named entity variations (e.g., “إبن الهيثم” vs. “ابن الهيثم”), and diacritics sensitivity (e.g., “المثانة” vs. “المثانة”). Based on these findings, Query2Doc [6] was

selected as the first technique because its pseudo-document generation mechanism directly addresses information poverty in short queries and vocabulary mismatch—the two dominant failure patterns identified. The complete rationale for this selection is documented in Section 3.4.

4.3 Query2Doc Results

This section presents the results of the initial Query2Doc experiments using Qwen 2.5 3B Instruct, evaluated on both dense and sparse retrieval.

4.3.1 Dense Retrieval

Table 4.5 presents the results of Query2Doc with dense retrieval.

Table 4.5: Query2Doc results with dense retrieval (mDPR). Improvement is computed relative to the dense baseline.

Metric	Baseline	Query2Doc	Improvement
NDCG@10	0.4993	0.5435	+8.9%
Recall@10	0.6156	0.6608	+7.3%
Recall@100	0.8407	0.8594	+2.2%
MRR	0.5328	0.5742	+7.8%

Query2Doc produced a substantial improvement of +8.9% in NDCG@10 over the dense baseline, representing the largest single-technique improvement achieved in the initial experiments. All metrics improved, with NDCG@10 and Recall@10 showing the largest gains. The smaller Recall@100 improvement (+2.2%) indicates that Query2Doc primarily improved *ranking* (moving relevant documents higher) rather than *discovery* (finding new relevant documents).

The query expansion statistics showed consistent enhancement: the median query length increased from 27 characters (original) to 250 characters (enhanced), representing an $8.45\times$ expansion ratio. The LLM-generated pseudo-documents provided additional vocabulary and semantic context that improved the mDPR encoder’s ability to match

queries with relevant passages.

The total runtime was approximately 40 minutes for 2,896 queries on a T4 GPU, achieved through the batch processing optimisations described in Section 3.4.4.

4.3.2 BM25 Retrieval

Table 4.6 presents the results of Query2Doc with BM25S sparse retrieval.

Table 4.6: Query2Doc results with BM25 retrieval (BM25S). Negative values indicate performance degradation.

Metric	Baseline	Query2Doc	Change
NDCG@10	0.4621	0.4090	−11.5%
Recall@10	0.5964	0.5384	−9.7%
Recall@100	0.8577	0.8155	−4.9%
MRR	0.4836	0.4342	−10.2%

In stark contrast to the dense results, Query2Doc *degraded* BM25 performance across all metrics, with NDCG@10 declining by 11.5%. This result, while negative, provided a valuable insight into the interaction between query enhancement techniques and retrieval paradigms.

4.3.2.1 Term Dilution Analysis

The performance degradation was attributed to **term dilution**: the BM25 scoring function (Equation 2.2) is sensitive to query length, as additional terms distribute the term-frequency weight across a larger vocabulary. When a short query such as “ما هي عاصمة السودان” (“What is the capital of Sudan,” 4 tokens) was expanded to 200+ tokens, the importance of the original keywords “عاصمة” (capital) and “السودان” (Sudan) was diluted in the BM25 scoring.

As noted in Section 3.4.6, this experiment did not implement the query repetition strategy recommended by Wang et al. [6], where the original query is repeated $n = 5$ times before concatenation. This omission was the primary cause of the BM25 degradation. The term dilution phenomenon and its resolution are further examined in the model

comparison experiments (Section 4.4.2), where models with Arabic-specific vocabulary were observed to partially overcome this limitation.

4.3.3 Comparison with the Original Query2Doc Paper

Table 4.7 compares the results of this implementation with those reported in the original Query2Doc paper [6].

Table 4.7: Comparison between the original Query2Doc results (English, MS-MARCO) and this implementation (Arabic, MIRACL).

Configuration	Original Paper	This Work
Language	English	Arabic
Dataset	MS-MARCO	MIRACL Arabic
LLM	GPT-3 (175B)	Qwen 2.5 3B
Prompting	Few-shot (4 examples)	Zero-shot
Dense NDCG@10 improvement	+2–5%	+8.9%
Sparse NDCG@10 improvement	+3–7%	−11.5%*

*Without query repetition; the original paper uses $n = 5$ repetition for sparse retrieval.

The dense retrieval improvement (+8.9%) exceeded the original paper’s reported range (+2–5%), despite using a model that is $58\times$ smaller and zero-shot rather than few-shot prompting. The mDPR baseline used here was not fine-tuned on MIRACL, unlike the strong retrievers used in the original paper, which likely contributed to the larger observed improvement.

4.4 Model Comparison Results

Following the methodology described in Section 3.5, ten open-source LLMs were evaluated using the identical Query2Doc pipeline. This section presents the results for both dense and sparse retrieval, including the analysis of dropped models.

4.4.1 Dense Retrieval Leaderboard

Table 4.8 presents the complete dense retrieval results, ranked by NDCG@10.

Table 4.8: Dense retrieval (mDPR) results for all models, ranked by NDCG@10. The improvement column is relative to the no-enhancement dense baseline.

Rank	Model	Params	NDCG@10	Recall@10	MRR	Δ NDCG
—	Baseline (no QE)	—	0.4993	0.6156	0.5328	—
1	Aya Expanse 8B	8B	0.6164	0.7256	0.6493	+23.5%
2	Jais-2 8B	8B	0.6018	0.7161	0.6356	+20.5%
3	Qwen3-8B	8B	0.5958	0.7119	0.6278	+19.3%
4	Qwen 2.5 7B	7B	0.5813	0.6952	0.6134	+16.4%
5	Qwen3-4B	4B	0.5691	0.6824	0.6015	+14.0%
6	Qwen 2.5 3B	3B	0.5435	0.6608	0.5742	+8.9%
7	Gemma 3 4B	4B	0.5391	0.6443	0.5761	+8.0%
8	Falcon-H1 3B	3B	0.5359	0.6484	0.5681	+7.3%
9	SILMA 2B	2B	0.5177	0.6289	0.5508	+3.7%
<i>Dropped model (see Section 4.4.3):</i>						
—	ALLaM 7B	7B	0.2550	0.3335	0.2708	−48.9%

All nine successfully evaluated models improved over the dense baseline, with improvements ranging from +3.7% (SILMA 2B) to +23.5% (Aya Expanse 8B). The top three models—Aya Expanse 8B, Jais-2 8B, and Qwen3-8B—all achieved improvements exceeding +19%, demonstrating that Query2Doc with appropriately selected LLMs can substantially enhance Arabic dense retrieval.




Figure 4.5: Dense retrieval NDCG@10 for all evaluated LLMs on the Query2Doc pipeline, sorted descending. The dashed reference line marks the mDPR no-QE baseline (0.499); each model uses its consistent thesis colour, matching Figure 4.8.




Figure 4.6: Grouped per-model bars for NDCG@10, Recall@10 and MRR on the Query2Doc dense pipeline; same model ordering as Figure 4.5.

4.4.2 BM25 Retrieval Results

Table 4.9 presents the BM25 retrieval results.

Table 4.9: BM25 retrieval (BM25S) results for all models, ranked by NDCG@10. The improvement column is relative to the no-enhancement BM25 baseline. Negative values indicate degradation.

Rank	Model	Params	NDCG@10	Recall@10	MRR	Δ NDCG
—	Baseline (no QE)	—	0.4621	0.5964	0.4836	—
1	Jais-2 8B	8B	0.5122	0.6448	0.5397	+10.8%
2	Aya Expanse 8B	8B	0.5046	0.6284	0.5377	+9.2%
3	Qwen 2.5 7B	7B	0.4682	0.6040	0.4905	+1.3%
4	Qwen3-8B	8B	0.4459	0.5806	0.4702	−3.5%
5	SILMA 2B	2B	0.4277	0.5550	0.4485	−7.4%
6	Qwen3-4B	4B	0.4145	0.5403	0.4415	−10.3%
7	Qwen 2.5 3B	3B	0.4090	0.5384	0.4342	−11.5%
8	Falcon-H1 3B	3B	0.4038	0.5311	0.4274	−12.6%
9	Gemma 3 4B	4B	0.3447	0.4532	0.3718	−25.4%

The BM25 results were markedly different from the dense results. Only three models improved over the BM25 baseline: Jais-2 8B (+10.8%), Aya Expanse 8B (+9.2%), and Qwen 2.5 7B (+1.3%). The remaining six models degraded BM25 performance, with Gemma 3 4B showing the largest decline (−25.4%).

The fact that two models—Jais-2 and Aya—achieved substantial BM25 improvements *without* implementing the query repetition strategy (Equation 3.2) is a notable finding. Jais-2’s success with BM25 is hypothesised to stem from its Arabic-centric vocabulary of 150,000 tokens (described in Section 2.4.1.2), which produces pseudo-documents with vocabulary that closely matches the BM25 index terms. Similarly, Aya Expanse’s multilingual training (23 languages with explicit Arabic support, Section 2.4.2.6) appears to generate lexically precise expansions.

4.4.3 Dropped Models Analysis

One model, ALLaM-7B, was dropped from the final comparison due to a severe performance issue. Following Dr. Tahani’s guidance that even negative results constitute valuable

research contributions, the case is analysed in detail.

4.4.3.1 ALLaM-7B

ALLaM-7B (described in Section 2.4.1.3) produced a catastrophic -48.9% decline in NDCG@10 (0.2550 vs. baseline 0.4993), performing *worse* than using no query enhancement at all. Investigation revealed a sentencepiece tokeniser artefact: the Unicode character “ ” (used internally to mark word boundaries) leaked into the decoded pseudo-documents. These special characters contaminated the query embeddings, causing the mDPR encoder to produce degraded representations.

This failure was attributed to ALLaM’s preview/alpha release status at the time of testing. The model had been accepted at ICLR 2025 [16] but was not yet in stable release, and the tokeniser behaviour had not been fully validated for generation tasks.

4.5 Key Findings and Analysis

This section synthesises the results across all experiments to identify the key findings and their implications.

4.5.1 Model Size Correlates with Dense Retrieval Performance

A positive association was observed between model parameter count and dense retrieval improvement, though it must be interpreted cautiously: across the full model set, parameter count is confounded with architecture, training-data volume, and tokeniser design, so the cross-model trend is only suggestive. A cleaner test comes from the Qwen family, whose four models share a common architecture and differ chiefly in scale and generation. Within this family, parameter count and dense NDCG@10 are strongly associated (Pearson $r = 0.94$; Spearman $\rho = 1.00$ across the 3B, 4B, 7B, and 8B models), supporting the size–quality relationship once architecture is held constant. Figure 4.7 illustrates the

broader relationship.




Figure 4.7: Relationship between model parameter count (billions) and gain over the mDPR baseline. A weak positive trend is visible across the full set, and the Qwen-family subset corroborates the size-quality correlation within a single architecture. ALLaM-7B is excluded as an outlier due to tokeniser artefacts.

Grouping models by size tier reveals a consistent pattern:

- **2–3B models:** +3.7% to +8.9% NDCG@10 improvement (SILMA, Qwen 2.5 3B, Falcon-H1)
- **4B models:** +8.0% to +14.0% (Gemma 3, Qwen3-4B)
- **7–8B models:** +16.4% to +23.5% (Qwen 2.5 7B, Qwen3-8B, Jais-2, Aya)

This suggests that larger models generate higher-quality pseudo-documents with richer vocabulary and more accurate factual content, leading to better-calibrated query representations in the dense embedding space.

4.5.2 Generational Improvement Within Model Families

The Qwen model family, with four models tested across two generations, provided a controlled comparison of generational improvements:

Table 4.10: Generational comparison within the Qwen model family.

Comparison	Qwen 2.5	Qwen3	Δ	Training Data
3–4B (Dense NDCG@10)	0.5435	0.5691	+4.7%	18T $\rightarrow$ 36T
7–8B (Dense NDCG@10)	0.5813	0.5958	+2.5%	18T $\rightarrow$ 36T

At both parameter scales, the newer Qwen3 generation outperformed Qwen 2.5, despite identical prompting and evaluation conditions. The improvement is attributed to Qwen3’s doubled training data volume (36T vs. 18T tokens), expanded language coverage (119 vs. 29 languages including 8 Arabic dialects), and wider attention architecture (described in Sections 2.4.2.3 and 2.4.2.4).

4.5.3 Arabic Specialisation vs. Multilingual Training

An important finding was that Arabic-specialised models did not consistently outperform multilingual models. Falcon-H1-3B, despite holding the highest Arabic OALL benchmark score ($\sim 62\%$) at the 3B scale, was outperformed on dense NDCG@10 by both Qwen 2.5 3B (a multilingual model with a lower OALL score) and Qwen3-4B. Similarly, Jais-2-8B, while achieving the best BM25 results, was surpassed on dense retrieval by Aya Expanse 8B, a model designed for 23 languages without Arabic-specific architecture.

This suggests that **a model’s Arabic-capability score on the Open Arabic LLM Leaderboard (OALL) does not directly predict its query expansion quality**. Falcon-H1-3B holds the highest OALL score at its parameter scale, yet it is outperformed on retrieval by multilingual models with lower OALL scores. The skills measured by OALL (reading comprehension, question answering, reasoning) differ from the ability to generate vocabulary-rich, lexically diverse pseudo-documents that improve retrieval; training-data volume and diversity appear to be more important factors for this specific task.

The notable exception is Jais-2’s BM25 performance: its custom 150,000-token Arabic vocabulary (Section 2.4.1.2) produces expansions with term distributions that align closely with BM25 index terms, making it uniquely effective for sparse retrieval.

4.5.4 Dense vs. BM25 Behaviour with Query Enhancement

The divergent response of dense and sparse retrieval to query enhancement is one of the most practically significant findings of this research. Table 4.11 summarises the pattern.

Table 4.11: Number of models improving over baseline, by retriever type.

Outcome	Dense (mDPR)	BM25S
Models improving	9 / 9	3 / 9
Models degrading	0 / 9	6 / 9
Avg improvement	+12.3%	−5.5%
Best improvement	+23.5%	+10.8%
Worst degradation	+3.7%	−25.4%

Dense retrieval benefited universally from query enhancement because the mDPR encoder operates in a continuous semantic space where additional context improves the query embedding representation. BM25, conversely, relies on exact term matching, where additional terms can dilute the weights of the original query keywords (Section 4.3.2.1).

This finding has direct implications for RAG system design: query enhancement techniques must be adapted to the retrieval paradigm, and a strategy that works well for dense retrieval should not be assumed to transfer to sparse retrieval without modification.

4.5.5 Model Selection Summary

Across the ten models evaluated, the following performance profiles emerged:

- a. **Strongest overall performance:** Aya Expanse 8B [21]. It achieved the highest dense NDCG@10 (0.6164, +23.5%) and the second-highest BM25 NDCG@10 (0.5046, +9.2%), the only model with strong improvements on *both* retrieval paradigms.
- b. **Strongest BM25 improvement:** Jais-2 8B [15]. It achieved the highest BM25 NDCG@10 (0.5122, +10.8%), benefiting from its Arabic-centric vocabulary for term-based matching.
- c. **Best efficiency profile:** Qwen3-4B [19]. At 4B parameters with no quantisation

required, it achieved +14.0% dense NDCG@10—nearly double the improvement of the 3B models—while fitting on a T4 GPU in FP16.

- d. **Temperature 0.1** was found to be optimal for deterministic pseudo-document generation, yielding +2.5% improvement over temperature 0.7 in controlled testing (Section 3.5.3).

4.6 BM25 Query Repetition Results

This section presents the results of the query repetition sweep described in Section 3.6. All nine models from the model comparison were evaluated across eight repetition configurations to determine whether the BM25 term dilution identified in Section 4.3.2.1 could be resolved without changing the underlying LLM.

Table 4.12 presents the full sweep results.

Table 4.12: BM25 nDCG@10 for nine models under query repetition. Columns $n \in \{1, 3, 5, 7, 10\}$ are fixed-count repetition (Query2Doc style); columns $\beta \in \{2, 4, 6\}$ are MuGI adaptive repetition where n is computed per-query as $n = \max(1, \lfloor |d| / (|q| \cdot \beta) \rfloor)$. MIRACL Arabic dev (2,896 queries). Best value per model in bold.

Model	$n=1$	$n=3$	$n=5$	$n=7$	$n=10$	$\beta=2$	$\beta=4$	$\beta=6$
BM25 alone (no QE)	0.4621	—	—	—	—	—	—	—
Aya Expanse 8B	0.5046	0.5652	0.5832	0.5849	0.5773	0.5855	0.5515	0.5256
Jais-2 8B	0.5122	0.5492	0.5529	0.5516	0.5436	0.5731	0.5521	0.5350
Qwen3-8B	0.4459	0.5181	0.5319	0.5328	0.5254	0.5254	0.4841	0.4591
Qwen 2.5 7B	0.4682	0.5294	0.5358	0.5331	0.5257	0.5320	0.4977	0.4774
Qwen3-4B	0.4145	0.5054	0.5239	0.5244	0.5188	0.5177	0.4678	0.4347
Gemma 3 4B	0.3447	0.4800	0.5178	0.5277	0.5239	0.4915	0.4184	0.3694
Qwen 2.5 3B	0.4090	0.5060	0.5185	0.5181	0.5116	0.5046	0.4551	0.4253
Falcon-H1 3B	0.4038	0.4881	0.5082	0.5112	0.5113	0.4979	0.4561	0.4266
SILMA 2B	0.4194	0.4783	0.4832	0.4829	0.4788	0.4494	0.4252	0.4203

Table 4.13 presents the supplementary Recall and MRR values at each model’s best configuration.

Five key observations emerged from the repetition sweep:

Table 4.13: BM25 retrieval metrics at each model’s best repetition configuration. Δ nDCG is the improvement from $n = 1$ (no repetition) to the best configuration.

Model	Best Config	nDCG@10	Recall@10	Recall@100	MRR	Δ nDCG
Aya Expanse 8B	$\beta = 2$	0.5855	0.7128	0.9300	0.6165	+0.0808
Jais-2 8B	$\beta = 2$	0.5731	0.7075	0.9217	0.6004	+0.0610
Qwen 2.5 7B	$n = 5$	0.5358	0.6765	0.9105	0.5586	+0.0675
Qwen3-8B	$n = 7$	0.5328	0.6695	0.9064	0.5590	+0.0868
Gemma 3 4B	$n = 7$	0.5277	0.6640	0.9002	0.5551	+0.1831
Qwen3-4B	$n = 7$	0.5244	0.6617	0.8980	0.5500	+0.1098
Qwen 2.5 3B	$n = 5$	0.5185	0.6501	0.9031	0.5494	+0.1095
Falcon-H1 3B	$n = 10$	0.5113	0.6456	0.8927	0.5379	+0.1074
SILMA 2B	$n = 5$	0.4832	0.6216	0.8747	0.5048	+0.0639

- a. **Query repetition brought all nine models above the BM25 baseline.** At $n = 1$ (no repetition), only three of nine models exceeded the BM25 baseline of 0.4621 (Aya at 0.5046, Jais-2 at 0.5122, and Qwen 2.5 7B at 0.4682). At the optimal repetition configuration, all nine models exceeded the baseline, including all six that had initially fallen below it.
- b. **Optimal repetition varied across models.** The 8B models (Aya, Jais-2, Qwen3-8B) performed best with MuGI adaptive repetition ($\beta = 2$), which scales repetition proportionally to pseudo-document length. The 3–4B models plateaued at fixed $n = 5$ –7. These differences likely reflect variation in pseudo-document length rather than model size per se.
- c. **Beyond the optimum, performance decreased symmetrically.** Excessive repetition over-weighted the original query tokens and suppressed the useful expansion vocabulary. The sharpest decline was visible in the $\beta = 4$ and $\beta = 6$ columns for the 3–4B models.
- d. **Aya Expanse 8B at $\beta = 2$ achieved 0.5855 nDCG@10**, improving BM25 by +26.7% over the baseline (0.4621) and recovering +0.0808 from the $n = 1$ result (0.5046). This confirmed that query repetition resolved the term-dilution degradation observed in the six models that initially underperformed the BM25 baseline.
- e. **Gemma 3 4B showed the largest absolute recovery ($\Delta = +0.1831$)**, rising from 0.3447 at $n = 1$ (the worst BM25 result of any model) to 0.5277 at $n = 7$. This

model had generated the longest pseudo-documents on average, causing the most severe term dilution at $n = 1$.




Figure 4.8: $n\text{DCG}@10$ on the BM25 retriever as a function of query-repetition count n for each LLM-generated pseudo-document. Six of nine models start below the BM25 baseline at $n=1$; all nine recover or exceed it at their optimum n . Colours match Figure 4.5.




Figure 4.9: Per-model Query2Doc gain on the Dense retriever versus the BM25 retriever (at $n=1$). Only Aya Expanse 8B and Jais-2 8B sit in the both-positive quadrant; most other models help Dense but degrade BM25 at $n=1$.

4.7 Hybrid Retrieval Fusion Results

This section presents the results of the hybrid fusion experiments described in Section 3.7. Table 4.14 reports the full Convex Combination (CC) α -sweep and both Reciprocal Rank Fusion (RRF) configurations.

Table 4.14: Hybrid retrieval fusion results on the MIRACL Arabic dev set (2,896 queries). CC = Convex Combination (min-max normalised per query, α controls Dense weight); RRF = Reciprocal Rank Fusion.

Method	nDCG@10	Recall@10	Recall@100	MRR
BM25S alone	0.4621	0.5964	0.8577	0.4836
mDPR alone	0.4993	0.6156	0.8407	0.5328
Hybrid CC $\alpha = 0.1$	0.5248	0.6412	0.9333	0.5584
Hybrid CC $\alpha = 0.2$	0.5533	0.6683	0.9413	0.5861
Hybrid CC $\alpha = 0.3$	0.5830	0.7014	0.9449	0.6161
Hybrid CC $\alpha = 0.4$	0.6137	0.7392	0.9454	0.6426
Hybrid CC $\alpha = 0.5$	0.6266	0.7478	0.9458	0.6577
Hybrid CC $\alpha = 0.6$	0.6051	0.7289	0.9440	0.6355
Hybrid CC $\alpha = 0.7$	0.5743	0.6963	0.9439	0.6049
Hybrid CC $\alpha = 0.8$	0.5384	0.6634	0.9416	0.5678
Hybrid CC $\alpha = 0.9$	0.4996	0.6297	0.9350	0.5257
Hybrid RRF $k = 20$	0.6267	0.7597	0.9467	0.6517
Hybrid RRF $k = 60$	0.6230	0.7553	0.9467	0.6490

Five key observations were drawn from the hybrid fusion results:

- Hybrid fusion achieved a +35.6% improvement over BM25 alone and +25.5% over mDPR alone (nDCG@10), substantially outperforming either retriever in isolation and confirming the complementarity identified in Section 4.1.1.
- RRF $k = 20$ (0.6267) and CC $\alpha = 0.5$ (0.6266) were numerically indistinguishable on nDCG@10. RRF had the edge on Recall@10 (+0.0119), while CC had the edge on MRR (+0.0060). Because RRF requires no hyperparameter tuning, it was adopted as the primary hybrid baseline for all subsequent experiments.
- The CC sweep was smooth and unimodal, peaking at $\alpha = 0.5$ and degrading symmetrically toward either extreme. At $\alpha = 0.9$ (almost pure Dense), the result

(0.4996) was essentially the mDPR-alone baseline (0.4993); at $\alpha = 0.1$ (almost pure BM25), the result (0.5248) exceeded BM25 alone; a possible explanation is that even a small Dense component contributed to document reranking. This confirmed that the two retrievers contributed roughly equally.

- d. RRF $k = 20$ slightly outperformed $k = 60$ (0.6267 vs. 0.6230 nDCG@10). The smaller k gives more weight to top-ranked items, which benefits nDCG@10 but does not affect Recall@100 (both 0.9467).
- e. The hybrid RRF $k = 20$ result of 0.6267 nDCG@10 established the non-QE hybrid baseline that all subsequent query enhancement methods were required to surpass.




Figure 4.10: Convex-combination α sweep for the no-QE hybrid baseline, showing nDCG@10, Recall@{10,100} and MRR. nDCG@10 peaks near $\alpha = 0.5$; Recall@100 is largely insensitive to α .

4.8 Corpus-Steered Query Expansion Results

This section presents the results of the Corpus-Steered Query Expansion experiments described in Section 3.8. The results are organised into two subsections: main CSQE results and the component ablation.

4.8.1 Main CSQE Results

Table 4.15 compares the CSQE-enhanced retrieval results against the relevant baselines.

Table 4.15: nDCG@10, Recall@10, Recall@100, and MRR for CSQE-enhanced retrieval on the MIRACL Arabic dev set. CSQE config: $k = 5$ first-pass, 2 corpus + 2 blind samples, $\alpha = 4$, Aya Expanses 8B.

Method	nDCG@10	Recall@10	Recall@100	MRR
BM25S alone	0.4621	0.5964	0.8577	0.4836
BM25+CSQE	0.6157	0.7447	0.9422	0.6380
mDPR alone	0.4993	0.6156	0.8407	0.5328
Dense+CSQE	0.5915	0.7073	0.8816	0.6225
Hybrid RRF $k = 20$ (no QE)	0.6267	0.7597	0.9467	0.6517

BM25+CSQE achieved 0.6157 nDCG@10, representing a +33.2% improvement over BM25 alone and, notably, nearly matching the no-QE hybrid baseline (0.6267) using only a single retriever. Dense+CSQE improved over mDPR alone by +18.5% but fell below BM25+CSQE. Dense+CSQE underperformed BM25+CSQE, indicating that the expanded query hurt rather than helped the Dense retriever in this configuration. This observation motivated the retriever-specific application experiments in Section 4.9.

4.8.2 Component Ablation

Table 4.16 presents the ablation results isolating the contributions of corpus-grounded and blind expansion components.

The key finding is that corpus-grounded and blind expansion components are complementary:

Table 4.16: BM25 retrieval metrics for corpus-only, blind-only, and combined CSQE expansion variants. All variants use $\alpha = 4$, Aya Expanse 8B, $k = 5$ first-pass.

Variant	nDCG@10	Recall@10	Recall@100	MRR
BM25 alone	0.4621	0.5964	0.8577	0.4836
013c: Corpus-only (4c+0b)	0.5381	0.6457	0.8790	0.5651
013d: Blind-only (0c+4b)	0.5752	0.7089	0.9201	0.6032
013: CSQE 2+2	0.6157	0.7447	0.9422	0.6380

the combined 2+2 system (0.6157) exceeded both blind-only (0.5752) and corpus-only (0.5381) variants individually, a +0.0405 improvement over the stronger single component. Blind-only outperformed corpus-only on BM25 because blind generation produces full answer paragraphs with diverse Arabic term variants, which benefits BM25’s vocabulary breadth scoring, while corpus-only extraction produces passage-level excerpts structurally similar to the original query. Corpus samples contribute by anchoring the expansion to attested Wikipedia vocabulary, reducing entity misidentification; blind samples diversify answer-space coverage beyond the first-pass retrieval window.

Table 4.17 presents the effect of the query repetition factor α on BM25+CSQE performance.

Table 4.17: Effect of query repetition factor α on CSQE BM25 retrieval nDCG@10.

α	BM25+CSQE nDCG@10
1	0.6095
2	0.6130
3	0.6154
4	0.6157

Performance improved monotonically from $\alpha = 1$ to $\alpha = 4$, but the total gain was minimal (+0.0062). At $\alpha = 1$, the system already captured 98.9% of the $\alpha = 4$ nDCG@10 result. The query weighting factor was therefore not a critical hyperparameter for this configuration.

4.9 CSQE with Hybrid Fusion

This section presents the results of combining CSQE with hybrid BM25–Dense fusion, following the retriever-specific application methodology described in Section 3.8.3.

4.9.1 Fusion Configuration Comparison

Table 4.18 presents the results for all three CSQE application strategies combined with both fusion methods.

Table 4.18: Results for three CSQE application strategies in hybrid fusion. BM25-expanded: BM25 receives CSQE query, Dense receives original query. Dense-expanded: BM25 receives original query, Dense receives CSQE query. Both-expanded: both retrievers receive CSQE query.

Strategy	Fusion	nDCG@10	Recall@10	Recall@20
Hybrid RRF $k = 20$ (no QE)	RRF	0.6267	0.7597	0.8101
Dense-expanded (BM25 raw + Dense+CSQE)	RRF	0.6474	0.7928	0.8302
Dense-expanded (BM25 raw + Dense+CSQE)	CC $\alpha = 0.4$	0.6588	0.7851	0.8249
Both-expanded (BM25+CSQE + Dense+CSQE)	RRF	0.6936	0.8290	0.8701
Both-expanded (BM25+CSQE + Dense+CSQE)	CC $\alpha = 0.5$	0.6959	0.8249	0.8681
BM25-expanded (BM25+CSQE + Dense raw)	CC $\alpha = 0.6$	0.7088	0.8302	0.8701
BM25-expanded (BM25+CSQE + Dense raw)	RRF	0.7137	0.8363	0.8701

BM25-expanded was the winning configuration despite giving the Dense retriever the original short query rather than the CSQE expansion. When CSQE is applied to both retrievers (Both-expanded, 0.6936), the two ranked lists become less divergent from each other—both are now based on expanded queries grounded in the same first-pass documents—which lowers the complementarity that RRF fusion relies upon. BM25-expanded preserves this complementarity: the BM25 run is enriched by CSQE vocabulary while the Dense run retains its independent semantic signal, making the two lists more complementary for fusion. A second, reinforcing factor concerns the Dense encoder itself: mDPR was fine-tuned on MS MARCO, whose queries are short web-search strings [?], whereas a CSQE-expanded query is very long (mean $\approx 1,500$ characters, against ≈ 30 for the original query). Feeding such a long, out-of-distribution input to the Dense encoder

degrades its ranking quality, so withholding the expansion from the Dense side both preserves list divergence and avoids this length penalty, while the lexical BM25 retriever still benefits from the added vocabulary. This finding illustrates a **retriever-specific query representation principle**: each retriever should receive the query format that maximises its divergence from the other retriever’s ranked list.

Table 4.19 presents the contribution of each CSQE component within BM25-expanded RRF.

Table 4.19: BM25-expanded RRF ablation: contribution of corpus and blind expansion components when fused with the original Dense run.

BM25 Expansion	BM25-expanded RRF nDCG@10
Corpus-only (4c+0b) + Dense orig	0.6616
Blind-only (0c+4b) + Dense orig	0.7082
CSQE 2+2 + Dense orig	0.7137

Table 4.20 presents the effect of the query repetition factor α on the best BM25-expanded RRF system.

Table 4.20: Effect of query repetition factor α on the best BM25-expanded RRF system. All values use the same CSQE 2c+2b expansion, varying only the α prepended to the BM25-side query.

α	BM25-expanded RRF nDCG@10
1	0.7123
2	0.7121
3	0.7130
4	0.7137

The α parameter was nearly flat for the final fused system: the $\alpha = 1$ result (0.7123) was only 0.0014 below $\alpha = 4$ (0.7137). This was consistent with Table 4.17 (BM25-alone α sweep): once CSQE expansion was combined with a strong dense run via RRF, the α parameter had negligible effect. The $\alpha = 4$ configuration was reported as the primary result because it matched the original CSQE experimental settings, but $\alpha = 1$ would be preferable for a production deployment (one fewer token per expanded query). Figure 4.11 summarises the full progression from baselines to the best system.




Figure 4.11: Progression of nDCG@10 from the BM25 baseline (0.462) through intermediate query-enhancement steps to the best system, BM25-only-expanded with RRF fusion, at 0.714. Annotations on each bar report the gain over the best previous configuration.

Table 4.21 summarises the nDCG@10 improvements of BM25-expanded RRF over all prior systems.

Table 4.21: nDCG@10 improvements of BM25-expanded RRF over prior systems.

Comparison	Δ nDCG@10	% change
BM25-expanded RRF vs BM25 alone	+0.2516	+54.5%
BM25-expanded RRF vs mDPR alone	+0.2144	+42.9%
BM25-expanded RRF vs best blind BM25 QE (Aya $\beta = 2$)	+0.1282	+21.9%
BM25-expanded RRF vs best blind Dense QE (Aya 8B)	+0.0973	+15.8%
BM25-expanded RRF vs Hybrid RRF (no QE)	+0.0870	+13.9%
Both-expanded vs BM25-expanded	-0.0178	Dense+CSQE hurts
Dense-expanded vs Hybrid RRF	+0.0207	Weakest config still beats

All three CSQE fusion configurations exceeded the no-QE hybrid baseline. The +0.0870 improvement of BM25-expanded over the hybrid baseline (+13.9%) isolated the contribution of CSQE beyond what hybrid fusion alone provides.

4.9.2 Overall System Progression

Table 4.22 summarises the nDCG@10 progression from the baseline through each stage of pipeline development.

Table 4.22: nDCG@10 progression from baseline to best system on the MIRACL Arabic dev set (2,896 queries). Δ cumulative measured from BM25S baseline.

Stage	System	nDCG@10	Δ cumulative
Baseline	BM25S alone	0.4621	—
Baseline	mDPR alone	0.4993	—
Blind QE	Best blind Dense (Aya 8B, $n = 1$)	0.6164	+23.5% over baseline
Blind QE	Best blind BM25 (Aya 8B, $\beta = 2$)	0.5855	+26.7% over baseline
Hybrid	Hybrid RRF $k = 20$ (no QE)	0.6267	+35.6% over baseline
CSQE alone	BM25+CSQE	0.6157	+33.2% over baseline
CSQE+Hybrid	BM25-expanded: BM25+CSQE + Dense RRF	0.7137	+54.5% over baseline

The progression demonstrated a consistent improvement through three stages of pipeline development. Blind QE improvements (Sections 4.3 and 4.6) confirmed that LLM-generated pseudo-documents are effective for Arabic information retrieval. The hybrid baseline (Section 4.7) established that lexical and semantic retrieval are complementary. CSQE (Section 4.8) introduced corpus grounding that outperformed blind QE on BM25. The final combination of corpus-steered expansion applied asymmetrically to hybrid retrieval (BM25-expanded RRF) achieved 0.7137 nDCG@10—a 54.5% improvement over BM25 alone and a 42.9% improvement over mDPR alone. Recall@100 = 0.9734 indicated that 97.3% of relevant documents appeared within the top 100 candidates, establishing a strong upper bound for downstream reranking.

4.10 Per-Query Error Analysis

This section presents the per-query error analysis described in Section 3.9, decomposing the CSQE improvement to identify where corpus-steered expansion helps and where it introduces regressions.

4.10.1 Win/Loss Distribution

Table 4.23 presents the per-query win/loss distribution.

Table 4.23: Per-query comparison of the CSQE+Hybrid system (BM25-expanded RRF) against the Aya blind BM25 $n = 1$ baseline (0.5046 nDCG@10) on 2,896 MIRACL Arabic dev queries. The system’s per-query-mean nDCG@10 is 0.6936; this is computed over individual queries and differs slightly from the corpus-level headline of 0.7137 reported in Table 4.22, which is produced by the official pooled evaluation.

Condition	Queries	%
CSQE+Hybrid > Blind	1,646	56.8%
Tie	770	26.6%
CSQE+Hybrid < Blind	480	16.6%
Mean delta (CSQE – Blind)	—	+0.1890

The CSQE+Hybrid system improved on 56.8% of queries and regressed on 16.6%, with a mean per-query delta of +0.1890 nDCG@10. The 26.6% of ties were predominantly queries where both systems achieved either 0.0 (no relevant document retrieved by either) or 1.0 (perfectly retrieved).

A total of 258 queries (8.9%) were classified as failures (CSQE+Hybrid nDCG@10 < 0.1). A corpus-membership check over the full 2,061,414-document MIRACL Arabic index confirmed that *every* one of these queries has at least one relevant passage present in the corpus: none are irretrievable, and there is no dataset-level ceiling or indexing gap. The failures are therefore genuine retrieval failures of two kinds. For 199 of the 258 queries the relevant passage is present but is missed by *all* methods, including the BM25 baseline—intrinsically hard queries for both lexical and semantic retrieval. The remaining 58 queries are retrievable by BM25 alone (43 of them with BM25 nDCG@10 ≥ 0.3) but are lost by the CSQE hybrid; these are genuine regressions in which expansion or fusion displaced a document that the bare query would otherwise have surfaced.




Figure 4.12: Distribution of the per-query difference $\Delta = \text{NDCG@10}(\text{CSQE+Hybrid}) - \text{NDCG@10}(\text{Aya-blind BM25})$ across all 2,896 MIRACL Arabic dev queries. Shaded bands mark the Big-Win region ($\Delta > 0.3$) and the strict Regression region ($\Delta < -0.1$).

4.10.2 First-Pass Quality as the Largest Modulator

Table 4.24 presents the performance split by BM25 first-pass retrieval quality.

Table 4.24: System performance split by whether the BM25 first-pass retrieved a relevant document at rank 1 (top-1, $\text{qrel} \geq 1$).

First-Pass Quality	n	Blind QE nDCG@10	CSQE+Hybrid nDCG@10	Δ
Relevant doc retrieved	1,061	0.6684	0.8877	+0.219
No relevant doc retrieved	1,835	0.4100	0.5814	+0.171
All queries	2,896	0.5046	0.6936	+0.189

Table 4.25 presents the performance decomposition by query length.

Table 4.25: System performance decomposed by the number of words in the original MIRACL query, using the 1–3 / 4–8 / 9+ word buckets adopted throughout this thesis.

Query Length	n	Blind QE	CSQE+Hybrid	Δ abs	Δ rel
Short (1–3 words)	147	0.369	0.530	+0.161	+43.6%
Medium (4–8 words)	2,495	0.506	0.703	+0.197	+38.8%
Long (9+ words)	254	0.567	0.698	+0.132	+23.3%

When first-pass retrieval succeeded (36.6% of queries), CSQE+Hybrid achieved 0.8877 nDCG@10—near-ceiling retrieval. The first-pass quality gap (+0.3063 between the first-pass success and failure groups) was larger than the overall system advantage over blind QE (+0.1890), confirming that first-pass retrieval quality is the largest single modulator of CSQE effectiveness. Figure 4.13 visualises this split. The improvement holds across all query lengths (Table 4.25): short queries (1–3 words, $n = 147$) gain the most *proportionally* (+43.6%), as they are the most underspecified and benefit from both vocabulary expansion and corpus grounding, while medium queries (4–8 words) gain the most in *absolute* terms (+0.197). Long queries (9+ words, $n = 254$) still gain (+0.132, +23.3%) but least, because their greater information content already mitigates the information-poverty problem that motivates query expansion. Figure 4.14 shows the per-bucket comparison.




Figure 4.13: CSQE+Hybrid versus Aya-blind BM25 baseline conditioned on whether the BM25 first-pass top-1 document is relevant ($\text{qrel} \geq 1$). The asymmetry between the relevant ($n = 1,061$) and not-relevant ($n = 1,835$) groups demonstrates first-pass quality as the largest behavioural modulator of CSQE.




Figure 4.14: Per-bucket comparison of Aya-blind BM25 and CSQE+Hybrid using the 1–3 / 4–8 / 9+ word buckets. CSQE+Hybrid improves every bucket; short queries gain the most proportionally (+43.6%) and medium queries the most absolutely (+0.197).

4.10.3 Big Wins: The Corpus Grounding Effect

All 1,061 big-win queries ($\Delta > 0.3$) shared a single pattern: CSQE $\text{nDCG@10} = 1.000$, Blind = 0.000. In each case, blind QE hallucinated a plausible but factually incorrect entity, while CSQE’s first-pass corpus document anchored the expansion to the correct Wikipedia article. Table 4.26 presents three representative examples.

This pattern validated the corpus grounding hypothesis: for niche or ambiguous Arabic terms, blind QE confidently generated a common modern interpretation, while corpus grounding locked onto the actual Wikipedia entity. The impact on query expansion was direct: the corpus anchor supplied the correct Arabic vocabulary (e.g. **مملوكي رباط قلاوون** for the first example, rather than the surgical terms produced by blind generation) that both BM25 and Dense scoring relied upon. Notably, for all three of these examples the bare query alone—with no expansion at all—already retrieved the correct article; the blind pseudo-document therefore actively *poisoned* a query the retriever had already solved, whereas the corpus-grounded expansion stayed on target and preserved the hit.

Table 4.26: Representative big-win queries illustrating the corpus grounding effect. In all three cases, CSQE achieved $\text{nDCG@10} = 1.000$ while the blind baseline achieved 0.000.

Query (gloss)	Blind generates	QE	CSQE to	grounds	CSQE
ما هو الرباط المنصوري؟ ("What is al-Ribat al-Mansuri?")	A surgical ligature/suture used in operations		A Mamluk-era Sufi lodge (<i>ribat</i>) endowed by Sultan al-Mansur Qalawun		1.000
ما هي الأسماء الخمسة؟ ("What are the Five Nouns?")	A list of popular given names (Muhammad, Adam, Ibrahim, ...)		The Arabic grammatical category of the Five Nouns (أب, ذو, فو, حم, أخ)		1.000
ما هو الفن الجزيري؟ ("What is Insular art?")	Modern environmental / land art		Insular (Hiberno-Saxon) art of the post-Roman British Isles		1.000

4.10.4 Regression Analysis

Table 4.27 presents the classification of the 367 regression queries.

Table 4.27: Classification of 367 regression queries ($\text{CSQE nDCG@10} < \text{Blind}$ by > 0.1).

Type	Count	%	Description
A: Strong BM25 hurt	191	52%	BM25 baseline ≥ 0.3 ; CSQE expansion introduces off-topic vocabulary
B: Poisoned first-pass	131	36%	BM25 baseline < 0.1 ; first-pass retrieves irrelevant doc $\rightarrow$ LLM grounding on wrong content
C: Partial BM25	45	12%	BM25 baseline 0.1–0.3; mixed quality




Figure 4.15: Type breakdown of the 367 queries that regress under CSQE+Hybrid ($\Delta < -0.1$). Type A (strong-BM25-hurt-by-expansion) accounts for 52%, Type B (poisoned first-pass) for 36%, and Type C (other) for 12%.

Type A: Strong BM25 Hurt by Expansion. Type A regressions (52%, 191 queries) occurred when BM25 already handled the query well (baseline $nDCG@10 \geq 0.3$) and the CSQE expansion introduced Arabic vocabulary from the corpus document that diluted the original query terms. This was the same term-dilution effect observed in Section 4.3.2.1, but in reverse: here the expansion itself (rather than the pseudo-document) reduced BM25 precision for queries that were already well handled.

Type B: Poisoned First-Pass. Type B regressions (36%, 131 queries) share a common mechanism: the query is short or ambiguous and its surface form collides with an Arabic homonym, personal name, or near-topic entity, so BM25’s first pass retrieves the *wrong* entity. CSQE then grounds its expansion on those off-topic documents and the final query misses—even when the model’s generation is otherwise clean—whereas blind expansion, which ignores the first pass, often answers correctly. The mode spans three sub-flavours. In a *tokenisation/dialect homonym* (e.g. “ماهو التطرف؟”, “what is extremism?”), four of the five first-pass documents match the token “ماهو” and surface a southern-dialect

article, a song, a novel, and an Iranian village rather than the extremism article, because “ماهو” collides with a dialectal phrase rather than the intended interrogative “ما هو”. In a *personal-name homonym* (e.g. “متی ولد نجیب محفوظ؟”), the first pass retrieves a different “نجیب باشا محفوظ”—a physician, not the novelist. In a *wrong-entity / near-topic collision* (e.g. “من هو مصمم موقع ویکیپدیا؟”), the first pass returns an article on Wikipedia’s ban in Turkey, leading the expansion to ground on an unrelated entity. In each case it is the first-pass error, not the generation, that drives the regression.

Implications. These two failure modes motivate two concrete design improvements: a *first-pass quality gate* that falls back to blind expansion when the top-1 BM25 document overlaps poorly with the query (addressing Type B), and *asymmetric expansion weighting* that down-weights the expansion when BM25 is already strong (addressing Type A). Both are developed as recommendations in Chapter 5.

4.11 Summary of All Experiments

Table 4.28 provides a consolidated view of all experiments conducted in this research, gathered here so that the full progression can be read at a glance after the individual results have been presented.

The experimental results demonstrate a progression from individual retriever baselines through blind query enhancement, hybrid fusion, and finally corpus-steered expansion combined with hybrid retrieval. The best system (BM25-expanded RRF: BM25+CSQE + Dense original query) achieved 0.7137 nDCG@10—a 54.5% improvement over the BM25 baseline and a 13.9% improvement over the no-QE hybrid system. The detailed results of each experiment were presented in Sections 4.6–4.10 above.

Table 4.28: Summary of all experiments grouped by phase: baselines, initial Query2Doc experiments, and expanded experiments covering query repetition, hybrid fusion, and corpus-steered query expansion. Bold values indicate the best result for each metric across all experiments.

Model / Method	Retriever	NDCG@10	Recall@10	MRR	Status
<i>Baselines</i>					
None (baseline)	Dense	0.4993	0.6156	0.5328	Baseline
None (baseline)	BM25	0.4621	0.5964	0.4836	Baseline
None (baseline)	BM25+Dense (RRF)	0.6267	0.7597	0.6517	Best no-QE
<i>Initial Query2Doc experiments</i>					
Qwen 2.5 3B	Dense	0.5435	0.6608	0.5742	OK
Qwen 2.5 3B	BM25	0.4090	0.5384	0.4342	Degraded
Falcon-H1 3B	Dense	0.5359	0.6484	0.5681	OK
Jais-2 8B	Dense	0.6018	0.7161	0.6356	Good
Jais-2 8B	BM25	0.5122	0.6448	0.5397	Best BM25
Qwen3-4B	Dense	0.5691	0.6824	0.6015	Good
ALLaM 7B	Dense	0.2550	0.3335	0.2708	Dropped
Aya 8B	Dense	0.6164	0.7256	0.6493	Best Dense (Qwen3)
Aya 8B	BM25	0.5046	0.6284	0.5377	Good
<i>Expanded experiments</i>					
Aya 8B ($\beta = 2$)	BM25	0.5855	0.7128	0.6165	Best BM25
Hybrid RRF $k = 20$	BM25+Dense	0.6267	0.7597	0.6517	Best no-QE
BM25+CSQE	BM25	0.6157	0.7447	0.6380	Good
Dense+CSQE	Dense	0.5915	0.7073	0.6225	Good
BM25-expanded RRF	BM25+Dense	0.7137	0.8363	0.7362	Best overall

Chapter 5

Conclusion and Recommendations

This chapter summarises the key findings of the research, discusses the challenges encountered during the experimental work, and presents recommendations for future research directions in Arabic query enhancement for Retrieval-Augmented Generation systems.

5.1 Conclusions

This thesis investigated the effectiveness of LLM-based query enhancement for Arabic information retrieval, focusing on the Query2Doc technique applied using small open-source models. The following conclusions were drawn from the experimental results.

Baseline establishment and error analysis. Dense (mDPR) and sparse (BM25S) retrieval baselines were successfully established on the MIRACL Arabic benchmark, confirming their complementary strengths: mDPR achieved superior ranking quality ($\text{NDCG@10} = 0.4993$), while BM25S achieved broader retrieval coverage ($\text{Recall@100} = 0.8577$). Error analysis on the dense baseline revealed that 34% of queries failed ($\text{NDCG@10} < 0.3$) and that the shortest queries underperformed, achieving only 72% of long-query performance ($\text{NDCG@10} = 0.345$ vs. 0.476). Although performance did not increase monotonically with length (the overall length–performance correlation was negligible, $r \approx -0.01$), the specific weakness of very short queries validated the selection of query expansion as the primary enhancement technique.

Query2Doc transfers effectively to Arabic. The Query2Doc technique, originally developed for English using GPT-3 (175B parameters) with few-shot prompting, was successfully adapted for Arabic zero-shot application using models as small as 3 billion

parameters. The initial experiment with Qwen 2.5 3B achieved a +8.9% improvement in NDCG@10, exceeding the original paper’s reported improvement range of +2–5% despite using a model 58 times smaller and without task-specific demonstrations.

Comprehensive model comparison. Ten open-source LLMs were evaluated using a standardised protocol across both dense and sparse retrieval. All nine viable models improved dense retrieval performance, with improvements ranging from +3.7% (SILMA Kashif-2B) to +23.5% (Aya Expanse 8B) in NDCG@10. Aya Expanse 8B was identified as the best overall model, achieving the highest dense improvement (+23.5%) and the second-highest BM25 improvement (+9.2%). Jais-2 8B was identified as the best model for sparse retrieval (+10.8% NDCG@10), uniquely benefiting from its 150,000-token Arabic-centric vocabulary. For resource-constrained environments, Qwen3-4B achieved +14.0% dense improvement while fitting on a T4 GPU in FP16 precision without quantisation.

Analytical findings on model characteristics. Several cross-cutting patterns were identified from the model comparison. Larger models tended to produce better dense-retrieval expansions, although across the heterogeneous model set parameter count is confounded with architecture and training data; the cleanest evidence comes from the Qwen family, where size and dense NDCG@10 are strongly associated once architecture is held constant (Spearman $\rho = 1.00$ across the 3B, 4B, 7B, and 8B models). Generational improvement within the Qwen family was also confirmed, with Qwen3 outperforming Qwen 2.5 at both the 3–4B and 7–8B scales. Notably, a model’s Arabic-capability score on the Open Arabic LLM Leaderboard (OALL) did not directly predict its query expansion quality; training-data volume and diversity were found to be more significant factors than Arabic-specific architecture or OALL ranking.

Dense and sparse retrieval respond differently to query enhancement. The most practically significant finding was the divergent behaviour of dense and sparse retrieval under query enhancement. Dense retrieval benefited universally (9 of 9 models improved), as additional semantic content improved the query embedding representation. BM25 was degraded by 6 of 9 models due to term dilution, where the pseudo-document vocabulary diluted the importance of original query keywords in the BM25 scoring function. Only models with strong Arabic vocabulary (Jais-2, Aya Expanse) or sufficient multilingual breadth (Qwen 2.5 7B) overcame this limitation. This finding has direct implications

for RAG system design: query enhancement strategies must be adapted to the retrieval paradigm, and dense retrieval systems are more robust recipients of LLM-generated query expansions.

Query repetition resolves sparse retrieval degradation. The BM25 term-dilution problem, initially identified for 6 of 9 models (Section 4.3.2.1), was fully resolved through query repetition. By prepending the original query β times before the pseudo-document, all nine models exceeded the BM25 baseline at their optimal β . Aya Expanses 8B with $\beta = 2$ achieved 0.5855 NDCG@10—a 26.7% improvement over the BM25 baseline and a substantial recovery from the 0.5046 result at $\beta = 1$ (Section 4.6).

Hybrid retrieval establishes a strong non-QE ceiling. Combining BM25S and mDPR through Reciprocal Rank Fusion (RRF $k = 20$) achieved 0.6267 NDCG@10—a 35.6% improvement over BM25 alone—without any query enhancement, confirming the complementarity of lexical and semantic retrieval observed in Section 4.1.1 (Section 4.7).

Corpus-steered expansion validates the corpus grounding hypothesis. The CSQE pipeline, which grounds LLM expansion in first-pass retrieved documents, achieved 0.6157 NDCG@10 on BM25 alone—nearly matching the no-QE hybrid baseline and substantially outperforming blind Query2Doc. Per-query error analysis confirmed the underlying mechanism: when first-pass retrieval succeeds, CSQE anchors the expansion to the correct Wikipedia entity, achieving near-ceiling performance (0.8877 NDCG@10 on first-pass-relevant queries). The combined system (BM25-expanded RRF: BM25+CSQE + Dense original query) achieved **0.7137 NDCG@10—a 54.5% improvement over BM25 alone and a 13.9% improvement over the no-QE hybrid baseline** (Sections 4.8 and 4.9).

Retriever-specific query representation is critical. A key finding is that applying CSQE asymmetrically—only to the BM25 retriever—outperformed applying it to both retrievers (BM25-expanded: 0.7137 > Both-expanded: 0.6959). Applying CSQE to both retrievers reduced the complementarity between their ranked lists, lowering the fusion ceiling; applying it only to BM25 preserved the Dense retriever’s independent semantic signal. This retriever-specific representation principle has practical implications for similar hybrid retrieval pipelines (Section 4.9).

Overall. LLM-based query enhancement is an effective, modular, and resource-efficient strategy for improving Arabic information retrieval. Small open-source models with 4–8 billion parameters can deliver substantial retrieval improvements without API costs or proprietary dependencies. Combined with corpus grounding and hybrid fusion, this approach raised retrieval quality from a BM25 baseline of 0.4621 to 0.7137 NDCG@10—a 54.5% improvement—making it a practical strategy for real-world Arabic RAG deployments.

5.2 Challenges

Several challenges were encountered during the course of this research.

1. **Resource constraints.** All experiments were conducted on Google Colab with NVIDIA T4 (15 GB) and A100 (40 GB) GPUs. This limited model sizes to 8 billion parameters (with 4-bit quantisation for larger models) and required significant engineering optimisation—batch processing, reduced token generation (128 tokens), and a two-notebook workflow separating query generation from retrieval evaluation—to achieve practical runtimes of approximately 40 minutes per full experiment.
2. **BM25 term dilution.** Query enhancement degraded BM25 performance for 6 of 9 models in the initial experiments. The query repetition strategy recommended by the original Query2Doc paper [6] ($n = 5$ repetitions of the original query before concatenation) was not implemented at that stage, as the focus was on establishing the baseline Query2Doc behaviour across models. This challenge was subsequently resolved through the query repetition technique (Section 4.6), which recovered BM25 performance for all six models that had initially degraded.
3. **Dropped model.** One of the ten models evaluated was unusable for Arabic query enhancement: ALLaM-7B, a preview-stage model, exhibited a sentencepiece tokeniser artefact that leaked special characters into the pseudo-documents, producing a catastrophic –48.9% decline in NDCG@10. This failure highlights the risk of relying on preview-stage models—whose tokeniser behaviour may not yet be validated for generation—in Arabic NLP tasks.

4. **Dataset scope.** The MIRACL Arabic benchmark contains exclusively Modern Standard Arabic text. The effectiveness of query enhancement for dialectal Arabic queries—Egyptian, Levantine, Gulf, Maghrebi, and other regional varieties—could not be evaluated within the scope of this work. Given the diglossic nature of Arabic, where users may formulate queries in dialect while knowledge bases are written in MSA, this represents a significant gap in the evaluation. A related limitation is that the MIRACL Arabic corpus provides no structural metadata—article boundaries, section headings, or passage hierarchy—beyond raw passage text, which prevented the knowledge-base-aware, “chunking-aware” query enhancement originally envisaged (Section 5.3).
5. **Single query enhancement technique.** Only the Query2Doc technique (generative query expansion via pseudo-document concatenation) was evaluated. Other approaches—HyDE (hypothetical document embedding), GRF (generative relevance feedback), query rewriting, and iterative refinement—were not compared. The results therefore reflect the potential of one specific technique rather than the broader landscape of query enhancement methods.
6. **Baseline retriever strength.** The mDPR dense retriever used in this work was pre-trained on MS MARCO and not fine-tuned on MIRACL Arabic. The magnitude of improvement observed may therefore differ when query enhancement is applied to retrieval models that have been specifically fine-tuned on Arabic data.
7. **First-pass quality dependence.** CSQE performance was found to be strongly conditioned on first-pass retrieval quality. When BM25 first-pass retrieval returned irrelevant documents (36% of regression cases), the LLM expansion was grounded on incorrect content, causing regressions that blind QE did not exhibit. Arabic homonym sensitivity in BM25 (e.g., retrieved as a dialectal phrase rather than a question prefix) was identified as the primary trigger (Section 4.10).

5.3 Recommendations for Future Work

Based on the findings and limitations of this research, the following recommendations are proposed for future work. These recommendations are ordered from direct extensions of the current work to broader research directions.

1. **Knowledge-base-aware query enhancement.** The current Query2Doc implementation generates pseudo-documents without any knowledge of the target corpus structure. Future work should investigate injecting structural information from the knowledge base—such as passage hierarchy, article boundaries, section headings, and metadata—into the query enhancement prompt. This “chunking-aware” approach would enable the LLM to generate context-aware expansions tailored to the actual content organisation of the knowledge base, potentially bridging the gap between query-side and corpus-side optimisation.
2. **Stronger embedding models.** The effectiveness of query enhancement should be evaluated with state-of-the-art Arabic-capable embedding models. BGE-M3 remains a strong, openly available multilingual retriever validated on MIRACL [12]; more recent encoders have since surpassed the older multilingual-E5-large, including the multilingual Qwen3-Embedding-8B [?] and the Arabic-centric Swan-Large [?], which outperforms multilingual-E5-large on most Arabic tasks. As these models are ranked primarily on general embedding benchmarks (MTEB/MMTEB and Arabic STS) rather than MIRACL Arabic retrieval specifically, a small validation run is advisable before adoption. This evaluation would establish whether query enhancement provides additive improvement even over strong retrievers, or whether the benefits diminish as the baseline retriever becomes more capable.
3. **Dialectal Arabic evaluation.** The effectiveness of query enhancement for dialectal Arabic queries should be evaluated using datasets that include regional Arabic varieties. The MIRACL Arabic benchmark contains exclusively Modern Standard Arabic text, so the extent to which these findings generalise to queries formulated in Egyptian, Levantine, Gulf, or Maghrebi Arabic remains an open question.

4. **Few-shot and chain-of-thought prompting.** The current work used exclusively zero-shot prompting for pseudo-document generation. Few-shot prompting with curated Arabic demonstration examples may improve pseudo-document quality by providing the model with explicit patterns for Arabic query expansion. Chain-of-thought reasoning before expansion—where the model analyses the query’s information needs before generating the pseudo-document—may further improve quality for complex or ambiguous queries.
5. **Multi-stage query enhancement.** The current approach applies a single round of query enhancement. Future work should investigate iterative refinement, where the results of a first retrieval pass inform a second round of query enhancement through a relevance feedback loop. This multi-stage approach could combine Query2Doc with re-ranking techniques, using the initially retrieved passages to generate more targeted query expansions.
6. **First-pass quality gate.** A quality filter should be implemented that checks lexical overlap between the top-1 BM25 document and the query before grounding the LLM expansion. When overlap falls below a threshold, the system should fall back to blind QE. This is expected to recover approximately 131 Type B regression queries identified in the error analysis (Section 4.10.4).
7. **Asymmetric expansion weighting.** For Type A regressions (strong BM25 queries), reducing the expansion weight α or using field-based BM25 indexing (original query in one field, expansion in another with lower boost) would prevent vocabulary dilution for queries that BM25 already handles well (Section 4.10.4).
8. **CSQE with stronger dense retrievers.** The current pipeline uses mDPR as a relatively weak dense baseline (Section 3.2.1). Testing CSQE+Hybrid with stronger dense retrievers such as BGE-M3 or Qwen3-Embedding-8B would assess whether the 54.5% BM25 improvement translates to similar gains over stronger dense retrievers, and whether the retriever-specific representation finding generalises.

Bibliography

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” arXiv.org, 2017. [Online]. Available: <https://arxiv.org/abs/1706.03762>
- [2] M. Song and M. Zheng, “A survey of query optimization in large language models,” arXiv.org, 2024. [Online]. Available: <https://arxiv.org/abs/2412.17558>
- [3] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive nlp tasks,” arXiv.org, 04 2021. [Online]. Available: <https://arxiv.org/abs/2005.11401>
- [4] S. Robertson and H. Zaragoza, “The probabilistic relevance framework: BM25 and beyond,” *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–389, 2009.
- [5] V. Karpukhin, B. Oğuz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W.-t. Yih, “Dense passage retrieval for open-domain question answering,” arXiv.org, 2020. [Online]. Available: <https://arxiv.org/abs/2004.04906>
- [6] L. Wang, N. Yang, and F. Wei, “Query2doc: Query expansion with large language models,” arXiv.org, 10 2023. [Online]. Available: <https://arxiv.org/abs/2303.07678>
- [7] S. Bruch, S. Gai, and A. Ingber, “An analysis of fusion functions for hybrid retrieval,” *ACM Transactions on Information Systems*, vol. 42, no. 1, 2024, arXiv:2210.11934.
- [8] L. Gao, X. Ma, J. Lin, and J. Callan, “Precise zero-shot dense retrieval without relevance labels,” 12 2022. [Online]. Available: <https://arxiv.org/pdf/2212.10496>

- [9] I. Mackie, S. Chatterjee, and J. Dalton, “Generative relevance feedback with large language models,” in *Proceedings of the 46th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2023.
- [10] X. Ma, Y. Gong, P. He, H. Zhao, and N. Duan, “Query rewriting for retrieval-augmented large language models,” arXiv.org, 10 2023. [Online]. Available: <https://arxiv.org/abs/2305.14283>
- [11] S. R. El-Beltagy and M. A. Abdallah, “Exploring retrieval augmented generation in arabic,” arXiv.org, 2024. [Online]. Available: <https://arxiv.org/abs/2408.07425>
- [12] J. Alsubhi, M. D. Alahmadi, A. Alhusayni, I. Aldailami, I. Hamdine, A. Shabana, Y. Iskandar, and S. Khayyat, “Optimizing rag pipelines for arabic: A systematic analysis of core components,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2506.06339v1>
- [13] Technology Innovation Institute, “Falcon-h1: A family of hybrid-head language models,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2507.22448>
- [14] A. Gu and T. Dao, “Mamba: Linear-time sequence modeling with selective state spaces,” arXiv.org, 2024. [Online]. Available: <https://arxiv.org/abs/2312.00752>
- [15] N. Sengupta, S. K. Sahu, B. Jia, S. Katipomu, H. Li, F. Krupber, W. Mahboob, W. Zuo, N. Luber, C. Axmed *et al.*, “Jais and jais-chat: Arabic-centric foundation and instruction-tuned open generative large language models,” arXiv.org, 2023. [Online]. Available: <https://arxiv.org/abs/2308.16149>
- [16] M. S. Bari, Y. Alnumay, N. A. Alzahrani, N. M. Alotaibi, H. A. Alyahya, S. AlRashed, F. A. Mirza, S. Z. Alsubaie, H. A. Alahmed, G. Alabduljabbar *et al.*, “ALLam: Large language models for arabic and english,” in *The Thirteenth International Conference on Learning Representations*, 2025. [Online]. Available: <https://openreview.net/forum?id=MscdsFVZrN>
- [17] SILMA AI, “SILMA kashif-2b-instruct: Arabic RAG-optimized language model,” 2024, model card, accessed March 2026. [Online]. Available: <https://huggingface.co/silma-ai/SILMA-Kashif-2B-Instruct-v1.0>

-
- [18] A. Yang, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Li, D. Liu, F. Huang, H. Wei *et al.*, “Qwen2.5 technical report,” arXiv.org, 2024. [Online]. Available: <https://arxiv.org/abs/2412.15115>
- [19] Qwen Team, “Qwen3 technical report,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2505.09388>
- [20] Gemma Team, Google DeepMind, “Gemma 3 technical report,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2503.19786>
- [21] J. Dang, V. Aryabumi, S. Dash, E. Beeching, S. Shandilya, A. Üstün, A. P. Singh, A. Bapna *et al.*, “Aya expanse: Connecting billions of people through language technology,” arXiv.org, 2024. [Online]. Available: <https://arxiv.org/abs/2501.01482>
- [22] OpenAI, “GPT-OSS,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2508.10925>
- [23] X. H. Lü, “BM25S: An efficient python implementation of BM25,” 2024, software library. [Online]. Available: <https://github.com/xhluca/bm25s>
- [24] I. Mackie, S. Chatterjee, and J. Dalton, “Generative and pseudo-relevant feedback for sparse, dense and learned sparse retrieval,” arXiv.org, 2023. [Online]. Available: <https://arxiv.org/abs/2305.07477>
- [25] X. Lei, Z. Lian, X. Zhang, and Y. Lin, “Corpus-steered query expansion with large language models,” arXiv.org, 2024. [Online]. Available: <https://arxiv.org/abs/2402.18031>
- [26] J. Zhang, M. Deng, Y. Zhou, S. Wang, and W. Li, “MUGI: Multi-granularity query expansion with llms for boosting sparse and dense retrieval,” in *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024. [Online]. Available: <https://aclanthology.org/2024.findings-emnlp.103>
- [27] Q. Zhang, F. Chen, Q. Yang, Y. Liu, and C. Wu, “Personalized query expansion with llms for dense and sparse retrieval,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2510.08935>

-
- [28] S. Yoon, S. Kim, and J. Park, “On the role of llms in query expansion for information retrieval,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2504.14175>
- [29] Y. Xia, J. Wang, Y. Zhang, and D. Li, “Knowledge-augmented retrieval: Bridging the gap between query and document representations,” 2025. [Online]. Available: <https://aclanthology.org/2025.naacl-long.216>
- [30] Z. Yang, X. Zhang, and Y. Huang, “Aligned query expansion for dense retrieval,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2507.11042>
- [31] X. Lei, Z. Lian, and X. Zhang, “ThinkQE: Chain-of-thought query expansion with large language models,” 2025. [Online]. Available: <https://aclanthology.org/2025.findings-emnlp.965>
- [32] O. Young, Y. Fan, R. Zhang, J. Guo, M. de Rijke, and X. Cheng, “GaQR: An efficient generation-augmented question rewriter,” in *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management (CIKM ’24)*, 2024, pp. 4228–4232.
- [33] C.-M. Chan, C. Xu, R. Yuan, H. Luo, W. Xue, Y. Guo, and J. Fu, “Rq-rag: Learning to refine queries for retrieval augmented generation,” arXiv.org, 03 2024. [Online]. Available: <https://arxiv.org/abs/2404.00610>
- [34] L. Louis and A. Bhusal, “Query augmentation for information retrieval (ir) using large language model (llm),” 2024.
- [35] K. Zhang, Z. Sun, W. Yu, X. Zang, K. Zheng, Y. Song, H. Li, and J. Xu, “Qe-rag: A robust retrieval-augmented generation benchmark for query entry errors,” arXiv.org, 2025. [Online]. Available: <https://arxiv.org/abs/2504.04062>
- [36] X. Zhang, N. Thakur, O. Ogundepo, E. Kamalloo, D. Alfonso-Hermelo, X. Li, Q. Liu, M. Rezagholizadeh, and J. Lin, “Miracl: A multilingual retrieval dataset covering 18 diverse languages,” arXiv.org, 2023. [Online]. Available: <https://arxiv.org/abs/2210.09984>
- [37] J. Lin, X. Ma, S.-C. Lin, J.-H. Yang, R. Pradeep, and R. Nogueira, “Pyserini: A python toolkit for reproducible information retrieval research

- with sparse and dense representations,” arXiv.org, 2021. [Online]. Available: <https://arxiv.org/abs/2102.10073>
- [38] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer, “QLoRA: Efficient finetuning of quantized language models,” arXiv.org, 2023. [Online]. Available: <https://arxiv.org/abs/2305.14314>